

CS 542 Class Challenge

Shuo Liu Ji Zhao
BUID: U28149952 BUID: U23540628
email: liushuo@bu.edu email: zhaoji@bu.edu

Task 1:

Based on the hint of task 1 we tried two different models VGG16, VGG19 for the classification as the pre-training model and it turned out that the accuracy of VGG16 is a little higher than VGG19 if we keep same other parameters and layers and finally we decided to use VGG16 and tuned certain hyperparameters in order to achieve the best accuracy. The deep neural network model VGG16 contains 16 weighted layers. The input of the network has a (224,224,3) shape and begins with two convolution layers, with a kernel size of (3x3), and the number of output channels is 64. Then a (2,2) max-pooling layer is followed to reduce the dimensionality. Two convolutional layers are applied again, with the (3x3) kernel size, this time the number of output channels is 128. As we go further into the network, the number of convolutional layers increases from two to three as a block, and the number of output channels increases from 128 to 256, and finally 512. Among these CNN blocks, (2,2) max-pooling is applied after every block. After the last pooling layer, we flatten the 512 output channels and two dense layers having 256 and 1 neuron were added with sigmoid acting on the last dense layer since we only have two classes. We tried adding Dropout layers as regulation after the flattening layer and before the output layer with a variety of rates from 0.1 to 0.25 and found this value did not cause high performance. Figure 1 is the architecture of our model:

Model: "model_8"			block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
Layer (type)	Output Shape	Param #	block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
input_7 (InputLayer)	[(None, 224, 224, 3)]	0	block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792	block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928	block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0	block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856	block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584	flatten_5 (Flatten)	(None, 25088)	0
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0	dense_feature (Dense)	(None, 256)	6422784
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168	dense_5 (Dense)	(None, 1)	257
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080	=====		
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080	Total params: 21,137,729		
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0	Trainable params: 6,423,041		
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160	Non-trainable params: 14,714,688		

Fig 1. The architecture of the task 1 model

When we compile the model, we combine the Focal Loss function (with k.binary_Crossentropy) in our model (alpha=0.25, gamma=2). A Focal Loss function addresses class imbalance during training in tasks like object detection. Focal loss applies a modulating term to the cross-entropy loss in order to focus learning on hard misclassified examples. It is a dynamically scaled cross-entropy loss, where the scaling factor decays to zero as confidence in the correct class increases. Intuitively, this scaling factor can automatically down-weight the contribution of easy examples during training and rapidly focus the model on hard examples. The optimizer we set is Adam, which involves a combination of two gradient descent methodologies: (1) Momentum: This algorithm is used to accelerate the gradient descent algorithm by taking into consideration the 'exponentially weighted average' of the gradients. Using averages makes the algorithm converge towards the minima at a faster pace; (2) Root mean square: Prop or RMSprop is an adaptive learning algorithm that tries to improve adaptive gradient. Instead of taking the cumulative sum of squared gradients, it takes the 'exponential moving average'.

Our learning rate was 0.0005 (which is better than 0.0001 in terms of accuracy). We also chose to use Epochs as 40. Based on figure 2, there is no overfitting so Early stopping is not needed. Note that when we increased the number of epochs to a much higher number, overfitting happened. Figure 2 is the accuracy and loss of our model over 40 epochs:

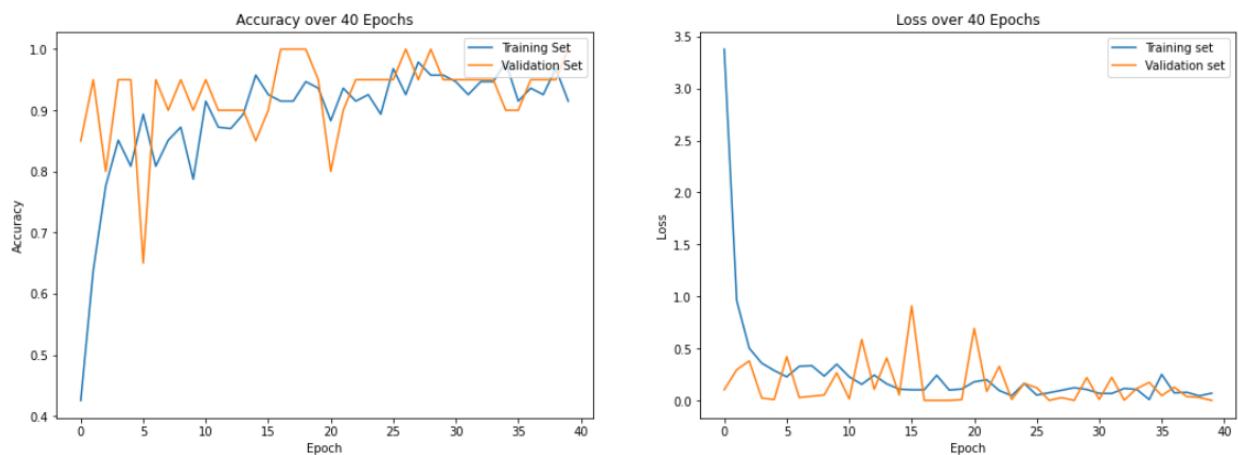


Fig 2. The accuracy and loss of our model over 40 epochs

For both the training and validation set accuracies, we can see that they are increasing at a sharp rate at the beginning while with a slow rate after around 4 epochs, but there are still some vibrations in between, especially for the validation set. This is most likely due to the minibatch training that is being used. This also explains why the validation accuracy is better than the training accuracy: during training, the model has access to all parameters so does validation but all parameters are only trained under part of the training set. The loss for the training and validation sets are all decreasing at a sharp rate at the beginning and after around 2 epochs the loss values are all less than 1, and for most epochs after 2, the loss values are less than 0.5. The usage of minibatch training might also be the reason that the training set loss is smoother than the validation set loss. To see the model's performance on test data, we evaluated the model on data from a test directory and received a more than 99% accuracy for

COVID19 Pneumonia and 99% accuracy for Normal, which means that the model has learned to generalize between two classes like shown in figure 3.

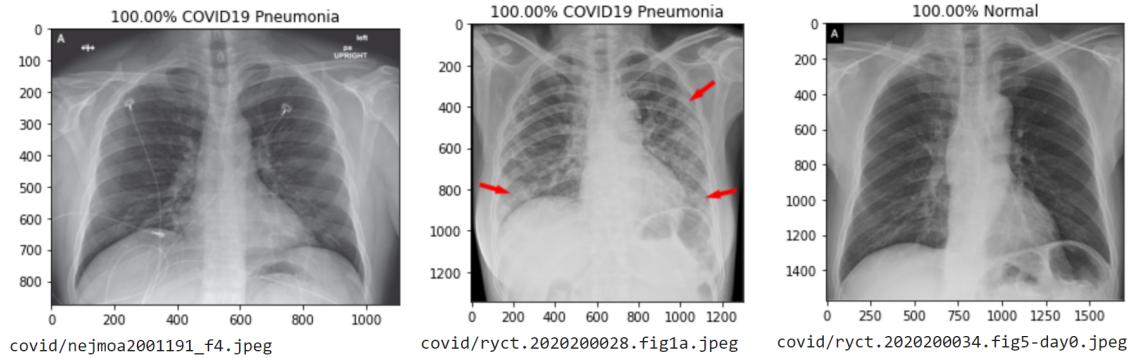


Fig 3. The accuracy for COVID19 and Normal

Finally, we show the visualization of the features of the test data in figure 4. The visualization plot was made using T-SNE to reduce the dimensionality of the data into a 2-dimensional feature space. The Blue data points are COVID-19 X-rays, and the orange data points are Normal X-rays. These features were evaluated from the input of the model to the “dense_feature” output. From this visualization, we can see besides the few Normal X-rays (several orange points) that are overlapping the COVID-19 X-rays (one blue point), the two classes are separable and this explains the high accuracy of our model: the model was able to create features that would be able to distinguish between the two classes.

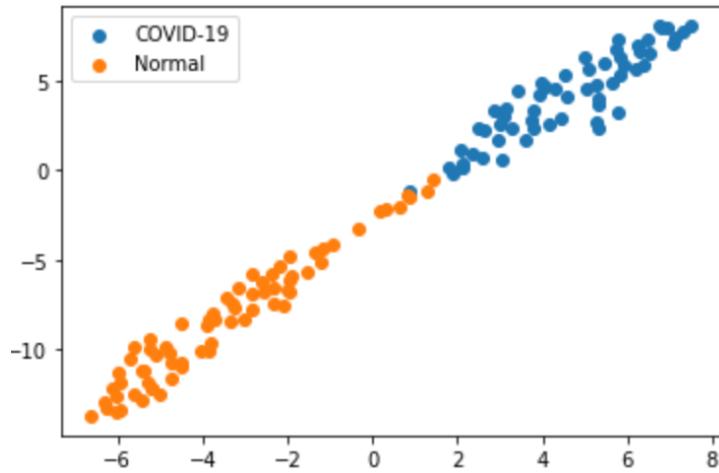


Fig 4. The visualization of COVID-19 X-rays and Normal X-rays

Task 2:

In this task, we were asked to train a deep neural model to classify an X-ray image into one of the following classes: normal, COVID-19, Pneumonia-Bacterial, and Pneumonia-Viral. VGG16 had also been applied in this task as well as an architecture called MobileNet for their outstanding image classification abilities.

Model 1:

Unlike the VGG16 model in Task 1, the output layer was changed to include 4 perceptrons, for 4 classes. For this task, the softmax function was used in the output layer and CategoricalCrossEntropy was applied as a loss function instead of the BinaryCrossEntropy loss in task 1. The feature extraction layer also contains 256 neurons. In this task, the epoch number was increased to 100 from 40 for this multi-class task. The optimizer we set is also Adam. It took about 15 mins to fit the model. The structure of the modified model in task 2 is as follows:

```
Model: "VGG16"
```

Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_14 (Flatten)	(None, 25088)	0
feature_dense (Dense)	(None, 256)	6422784
dense_14 (Dense)	(None, 4)	1028
<hr/>		
Total params: 21,138,500		
Trainable params: 6,423,812		
Non-trainable params: 14,714,688		

Fig 5. The structure and parameters of the modified VGG16 model

In addition to the VGG16 model, we applied MobileNet in this task.

Model 2:

MobileNet, as shown in the below figure, has a smaller structure, with less computation and higher precision that can be used on mobile terminals and embedded devices. Based on a deep separable roll product, Mobilenet uses two global hyperparameters to maintain a balance between efficiency and accuracy. The core idea of MobileNet is the decomposition of the convolution kernel. Standard convolutions can be decomposed into depthwise convolutions and pointwise convolutions with convolution kernels by using depthwise separable convolutions. A depthwise convolution filter convolved each channel, and the convolution was used to combine the outputs of the depthwise convolutional layers.

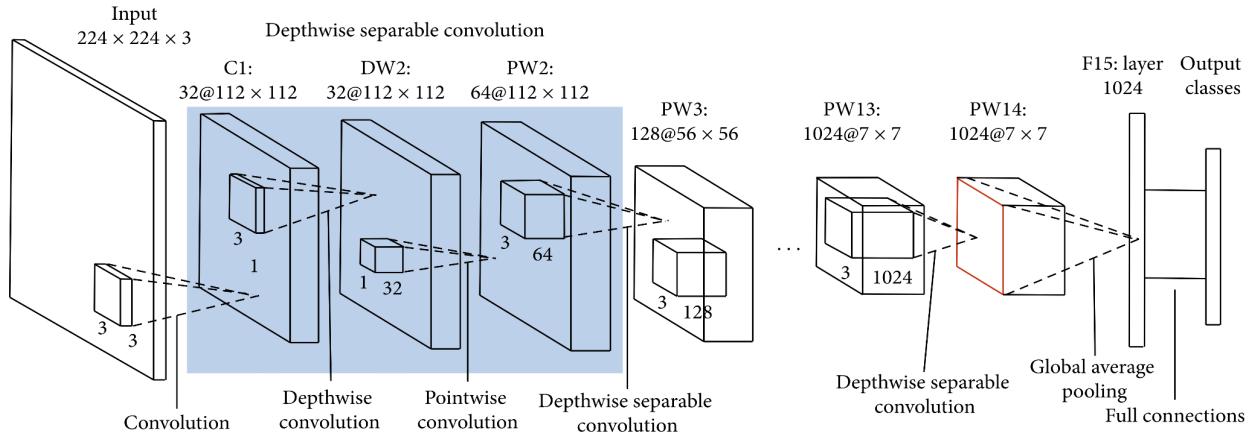


Fig 6. The structure of the MobileNetVI

The architecture of MobileNet has the same Input as the VGG16 architecture, with the image of the X-rays (224,224,3), but the rest of the architecture is very different in comparison.

The first layer of this MobileNet is a (3x3) standard convolution, followed by the accumulation of Depthwise convolutions. As can be seen, some depthwise convolutions will continue to downsample with strides=2. Then, the features are changed to (1x1) by average pooling, a fully connected layer is added according to the predicted class size, and the last Softmax layer. The entire network has 28 layers (excluding Average-Pool and Softmax). For parameters, it mainly focuses on (1x1) convolutions. In addition, the fully connected layer accounts for part of the parameters. Instead of using the last three layers of the standard model, we use a flatten and three fully connected layers to extract features and get the output. We also added a global-average pooling layer in the model which preserved the spatial information extracted by the previous convolutional layers and pooling layers. The modified model structure in Task 2 is as follows:

Model: "MobileNet"		
Layer (type)	Output Shape	Param #
mobilenet_1.00_224 (Functional)	(None, 7, 7, 1024)	3228864
global_average_pooling2d_5 (GlobalAveragePooling2D)	(None, 1024)	0
flatten_11 (Flatten)	(None, 1024)	0
dense_9 (Dense)	(None, 512)	524800
feature_dense (Dense)	(None, 256)	131328
dense_10 (Dense)	(None, 4)	1028

Total params: 3,886,020
Trainable params: 3,864,132
Non-trainable params: 21,888

Fig 7. The structure and parameters of the modified MobileNet model

The number of channels increases after each Dense block, going from 32 to 64, to 128, to 512, and finally 1024. Unlike VGG16, which requires 6.4 million parameters to train, the MobileNet architecture requires fewer parameters to train, which means less time cost. Batch Normalization exists after almost every convolutional layer in MobileNet to speed up training convergence and improve accuracy.

Accuracy:

The learning rate in task 2 for both models was 0.0001 because our model's accuracy dropped when we tried a learning rate greater than 0.0001. When we set the learning rate to be less than 0.0001, the training process became too slow and not efficient enough.

We tried many methods to avoid overfitting. Firstly, we used the early-stop method, but every training session stopped early, we would always have a model that would underfit. Secondly, we tried adding dropout layers to our models, but they performed worse than models without them. Moreover, as is known to us, a higher epoch might lead to a worse accuracy because of overfitting. Therefore, we tried to change the number of epochs to increase our models' accuracy.

Accuracy and Loss of Model 1: VGG16;

Parameters: Epoch = 100; Activation function = Softmax;

For this model, we used 100 epochs because when we tried higher epochs, the training and validation set began to overfit, and the performance of the model was significantly worse on the test set. When we used fewer epochs, there would be underfitting and the model performed worse or had no change on the test set too.

In order to get higher accuracy, we compared the accuracies of VGG16 and VGG19. We finally chose VGG16 because of its better performance even though VGG19 is more complex than VGG16.

Here is the accuracy of VGG16:

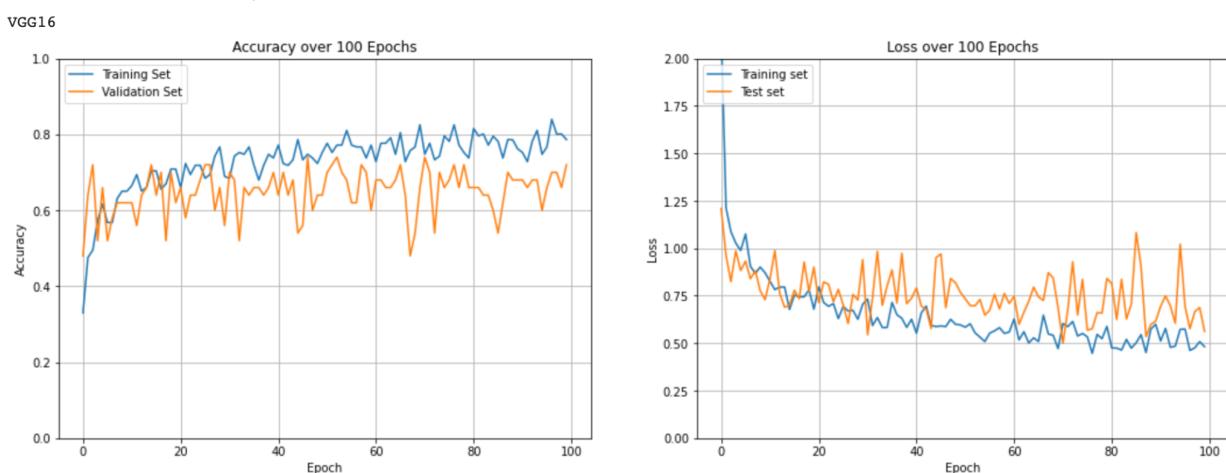


Fig 8. The accuracy and loss of VGG16

In task 1, the training accuracies were always lower than the validation accuracies because of the small amount of validation data. However, in task 2, the training accuracies were higher than the ones in the validation set. Around epoch 65 we could see that the model started to overfit the training set. From the loss graph on the right side of figure 8, there is a peak around epoch 65 which is consistent with the accuracy graph on the right. After training VGG16, it could be overfitting. To possibly improve performance, we can add more layers to the network so that it can constantly learn new nonlinear features and thus be able to classify.

Accuracy and Loss of Model 2: MobileNet;

Parameters: Epoch = 100; Activation function = Softmax;

For this model, we also used 100 epochs. We finally chose MobileNetV1 among MobileNetV1, V2, and V3 because of its better performance.

Here is the accuracy of MobileNet:

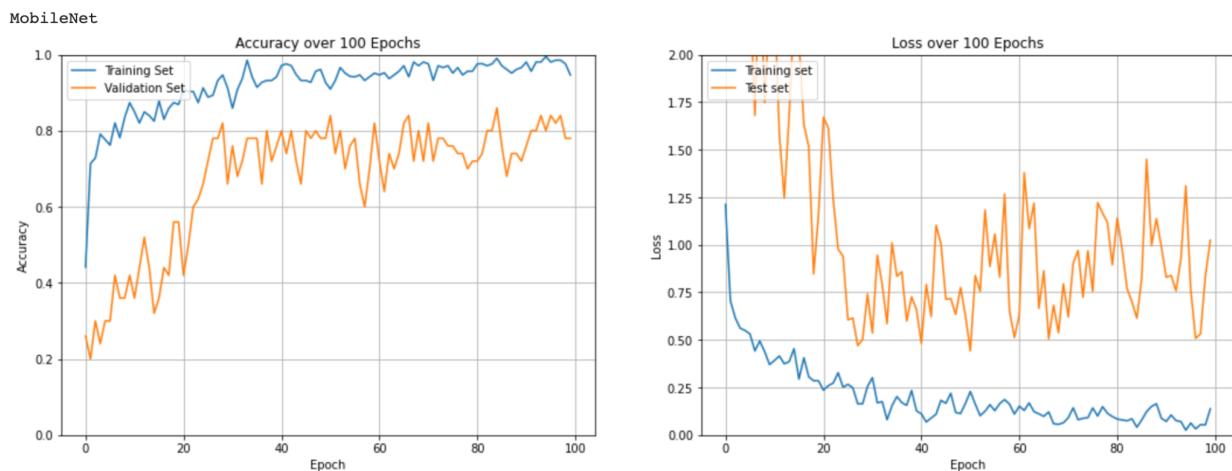


Fig 9. The accuracy and loss of MobileNet

Unlike VGG16, the trajectory of the accuracy of the training and validation improves over the number of epochs. This model performed well on the training set with really high accuracies, but the fluctuation of accuracy was too large for the validation set. Although the model didn't perform very well graphically, even after we added the dropout layer, the accuracy fluctuated greatly for the validation set and decreased for the test set. After repeated comparison and modification of model parameters, the current model performed best on the test set with an accuracy of about 0.8333. It performed better than VGG16 in task 2 with an accuracy of about 0.7500, which is shown in table 1.

Model	Epoch	#Dropout Layer	Activation function	Test Accuracy
VGG16	100	0	Softmax	0.7500
MobileNet	100	0	Softmax	0.8333

Table 1. Some information on models

T-SNE of two Models: VGG16 and MobileNet;

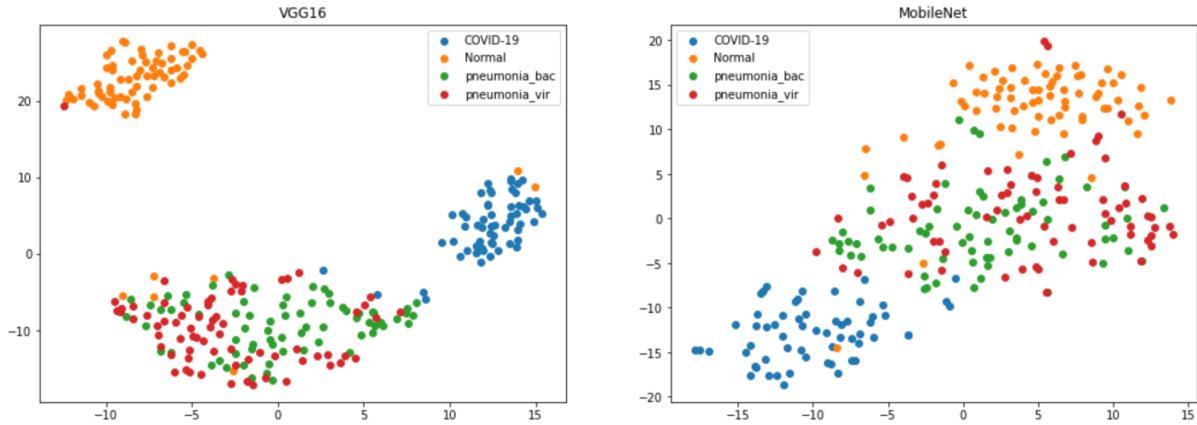


Fig 10. T-SNE of VGG16 (left) and MobileNet (right)

The visualizations of the features learned by the two models are shown in figure 10. Like in task 1, we also used T-SNE to reduce the dimensionality to two models to visualize the test data. In the two plots above, the blue data points are COVID-19 X-rays, orange data points are normal X-rays, green data points are pneumonia-bacterial X-rays, and red represents pneumonia-viral X-rays. The left side is the modified VGG16 model while the right side shows the modified MobileNet model. For the VGG16, These features were evaluated from the input of the model to the “feature_dense” output. We can see that in this model, features extracted from the test set were able to clearly separate normal and COVID-19 X-rays, though there is a small overlap of normal by pneumonia-viral data points and a small overlap of COVID-19 by pneumonia-bacterial data points. We speculate that few normal people have similar cinematographic signs to pneumonia-viral patients, and so do COVID-19 patients with normal people. We also see that the VGG16 model lacks the ability to distinguish the Pneumonia-bacterial and Pneumonia-viral X-rays, which may be due to the prime similarity of cinematographic signs with minor details that the model can not learn. Compared to the VGG16 model, the MobileNet model shows obviously better performance in terms of distinguishing the normal and COVID-19 from Pneumonia-bacterial and Pneumonia-viral X-rays for fewer overlap regions. It seems that for both models it is not difficult to separate X-rays of COVID-19 patients and normal people. We might select Pneumonia-bacterial and Pneumonia-viral X-rays for a binary classification case and try to use transfer learning for more clear feature extraction.

Compared to the T-SNE visualization in task 1, the modified VGG16 model in task 2 has the ability to learn how different the normal X-rays and COVID-19 X-rays are due to the long distance between two clusters. Nonetheless, the VGG16 lacks the ability to extract features to clearly separate Pneumonia-bacterial and Pneumonia-viral X-rays clusters.

Class Challenge: Image Classification of COVID-19 X-rays

Task 1 [Total points: 30]

Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib  
conda install numpy  
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib  
pip install numpy  
pip install sklearn
```

Data

Please download the data using the following link: [COVID-19](https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view) (https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view).

- After downloading 'Covid_Data_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
|--all  
|-----train  
|-----test  
|--two  
|-----train  
|-----test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

In [4]:

```
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

In []:

```
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))
```

Found GPU at: /device:GPU:0

[20 points] Binary Classification: COVID-19 vs. Normal

In []:

```
import os
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

os.environ['OMP_NUM_THREADS'] = '1'
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'
tf.__version__
```

Out[]:

'2.8.0'

Load Image Data

In []:

```
DATA_LIST = os.listdir('/content/drive/MyDrive/challenge/two/train')
DATASET_PATH = '/content/drive/MyDrive/challenge/two/train'
TEST_DIR = '/content/drive/MyDrive/challenge/two/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU
               # runs out of memory
NUM_EPOCHS = 40
LEARNING_RATE = 0.0005 # start off with high rate first 0.001 and experiment with
                       # reducing it gradually
```

Generate Training and Validation Batches

In []:

```
train_datagen = ImageDataGenerator(rescale=1./255, rotation_range=50, featurewise_center = True,
                                    featurewise_std_normalization = True, width_shift_range=0.2,
                                    height_shift_range=0.2, shear_range=0.25, zoom_range=0.1,
                                    zca_whitening = True, channel_shift_range = 20,
                                    horizontal_flip = True, vertical_flip = True,
                                    validation_split = 0.2, fill_mode='constant')

train_batches = train_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                 shuffle=True, batch_size=BATCH_SIZE,
                                                 subset = "training", seed=42,
                                                 class_mode="binary")

valid_batches = train_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                 shuffle=True, batch_size=BATCH_SIZE,
                                                 subset = "validation", seed=42,
                                                 class_mode="binary")
```

Found 104 images belonging to 2 classes.

Found 26 images belonging to 2 classes.

```
/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/image_data_generator.py:342: UserWarning: This ImageDataGenerator specifies `zca_whitening` which overrides setting of `featurewise_std_normalization`.

warnings.warn('This ImageDataGenerator specifies '
```

[10 points] Build Model

Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

In []:

```
from tensorflow.keras import Sequential, layers
from keras import models
from keras.models import *
from keras.layers import *
from keras.preprocessing.image import *
from keras.utils import *
from keras.optimizers import *
from keras.applications import *
from keras.applications import imagenet_utils
from keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint, LearningRateScheduler
from keras.applications import vgg16

# 模型 VGG161.0

base_model = vgg16.VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False
x = base_model.output
# x = GlobalAveragePooling2D()(x)
# x = Dense(128, activation="relu")(x)
# x = Dropout(0.5)(x)
x = layers.Flatten()(x)
# x = GlobalAveragePooling2D()(x)

x = Dense(256, name='dense_feature')(x)

out = Dense(1, activation="sigmoid")(x)
model = Model(base_model.input, out)

# model = models.Sequential()
# model.add(base_model)
# model.add(layers.Flatten())
# model.add(layers.Dense(256, activation='relu', name='dense_feature'))
# model.add(layers.Dropout(0.1))
# model.add(layers.Dense(1, activation='sigmoid'))

# model.build(input_shape=(224, 224, 3))
model.summary()
```

Model: "model_9"

Layer (type)	Output Shape	Param #
<hr/>		
input_3 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_2 (Flatten)	(None, 25088)	0
dense_feature (Dense)	(None, 256)	6422784
dense_2 (Dense)	(None, 1)	257
<hr/>		

Total params: 21,137,729

Trainable params: 6,423,041

Non-trainable params: 14,714,688

In []:

```
import keras.backend as K
from keras import optimizers

# focal loss
def focal_loss(alpha=0.25, gamma=2.0):
    def focal_crossentropy(y_true, y_pred):
        bce = K.binary_crossentropy(y_true, y_pred)

        y_pred = K.clip(y_pred, K.epsilon(), 1.- K.epsilon())
        p_t = (y_true*y_pred) + ((1-y_true)*(1-y_pred))

        alpha_factor = 1
        modulating_factor = 1

        alpha_factor = y_true*alpha + ((1-alpha)*(1-y_true))
        modulating_factor = K.pow((1-p_t), gamma)

        # compute the final loss and return
        return K.mean(alpha_factor*modulating_factor*bce, axis=-1)
    return focal_crossentropy

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0005), loss=focal_loss(), metrics=["accuracy"])
```

[5 points] Train Model

In []:

```
#FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN=train_batches.n//train_batches.batch_size
STEP_SIZE_VALID=valid_batches.n//valid_batches.batch_size

history = model.fit_generator(train_batches,
                               steps_per_epoch=STEP_SIZE_TRAIN,
                               epochs = 40,
                               validation_data=valid_batches,
                               validation_steps=STEP_SIZE_VALID,
                               shuffle=True)
# raise NotImplementedError("Use the model.fit function to train your network")
```

```
Epoch 1/40
10/10 [=====] - 7s 632ms/step - loss: 3.377
6 - accuracy: 0.4255 - val_loss: 0.1022 - val_accuracy: 0.8500
Epoch 2/40
10/10 [=====] - 5s 532ms/step - loss: 0.958
4 - accuracy: 0.6383 - val_loss: 0.2949 - val_accuracy: 0.9500
Epoch 3/40
10/10 [=====] - 5s 550ms/step - loss: 0.501
9 - accuracy: 0.7766 - val_loss: 0.3791 - val_accuracy: 0.8000
Epoch 4/40
10/10 [=====] - 6s 551ms/step - loss: 0.357
8 - accuracy: 0.8511 - val_loss: 0.0221 - val_accuracy: 0.9500
Epoch 5/40
10/10 [=====] - 6s 575ms/step - loss: 0.286
9 - accuracy: 0.8085 - val_loss: 0.0081 - val_accuracy: 0.9500
Epoch 6/40
10/10 [=====] - 6s 562ms/step - loss: 0.226
7 - accuracy: 0.8936 - val_loss: 0.4203 - val_accuracy: 0.6500
Epoch 7/40
10/10 [=====] - 5s 557ms/step - loss: 0.327
3 - accuracy: 0.8085 - val_loss: 0.0273 - val_accuracy: 0.9500
Epoch 8/40
10/10 [=====] - 6s 537ms/step - loss: 0.334
0 - accuracy: 0.8511 - val_loss: 0.0398 - val_accuracy: 0.9000
Epoch 9/40
10/10 [=====] - 6s 560ms/step - loss: 0.233
7 - accuracy: 0.8723 - val_loss: 0.0519 - val_accuracy: 0.9500
Epoch 10/40
10/10 [=====] - 6s 564ms/step - loss: 0.347
9 - accuracy: 0.7872 - val_loss: 0.2649 - val_accuracy: 0.9000
Epoch 11/40
10/10 [=====] - 5s 561ms/step - loss: 0.226
4 - accuracy: 0.9149 - val_loss: 0.0138 - val_accuracy: 0.9500
Epoch 12/40
10/10 [=====] - 6s 530ms/step - loss: 0.154
8 - accuracy: 0.8723 - val_loss: 0.5866 - val_accuracy: 0.9000
Epoch 13/40
10/10 [=====] - 6s 580ms/step - loss: 0.242
4 - accuracy: 0.8700 - val_loss: 0.1068 - val_accuracy: 0.9000
Epoch 14/40
10/10 [=====] - 6s 566ms/step - loss: 0.158
8 - accuracy: 0.8936 - val_loss: 0.4080 - val_accuracy: 0.9000
Epoch 15/40
10/10 [=====] - 6s 619ms/step - loss: 0.107
9 - accuracy: 0.9574 - val_loss: 0.0516 - val_accuracy: 0.8500
Epoch 16/40
10/10 [=====] - 6s 566ms/step - loss: 0.101
7 - accuracy: 0.9255 - val_loss: 0.9081 - val_accuracy: 0.9000
Epoch 17/40
10/10 [=====] - 5s 552ms/step - loss: 0.101
6 - accuracy: 0.9149 - val_loss: 1.6950e-04 - val_accuracy: 1.0000
Epoch 18/40
10/10 [=====] - 5s 538ms/step - loss: 0.241
4 - accuracy: 0.9149 - val_loss: 3.4810e-05 - val_accuracy: 1.0000
Epoch 19/40
10/10 [=====] - 5s 570ms/step - loss: 0.098
9 - accuracy: 0.9468 - val_loss: 1.6701e-04 - val_accuracy: 1.0000
Epoch 20/40
10/10 [=====] - 5s 527ms/step - loss: 0.109
4 - accuracy: 0.9362 - val_loss: 0.0066 - val_accuracy: 0.9500
Epoch 21/40
```

```
10/10 [=====] - 5s 533ms/step - loss: 0.178
6 - accuracy: 0.8830 - val_loss: 0.6926 - val_accuracy: 0.8000
Epoch 22/40
10/10 [=====] - 6s 557ms/step - loss: 0.198
3 - accuracy: 0.9362 - val_loss: 0.0858 - val_accuracy: 0.9000
Epoch 23/40
10/10 [=====] - 6s 557ms/step - loss: 0.094
8 - accuracy: 0.9149 - val_loss: 0.3274 - val_accuracy: 0.9500
Epoch 24/40
10/10 [=====] - 5s 552ms/step - loss: 0.045
0 - accuracy: 0.9255 - val_loss: 0.0078 - val_accuracy: 0.9500
Epoch 25/40
10/10 [=====] - 6s 590ms/step - loss: 0.162
8 - accuracy: 0.8936 - val_loss: 0.1629 - val_accuracy: 0.9500
Epoch 26/40
10/10 [=====] - 5s 565ms/step - loss: 0.050
7 - accuracy: 0.9681 - val_loss: 0.1214 - val_accuracy: 0.9500
Epoch 27/40
10/10 [=====] - 6s 562ms/step - loss: 0.073
9 - accuracy: 0.9255 - val_loss: 7.2357e-07 - val_accuracy: 1.0000
Epoch 28/40
10/10 [=====] - 5s 547ms/step - loss: 0.097
1 - accuracy: 0.9787 - val_loss: 0.0262 - val_accuracy: 0.9500
Epoch 29/40
10/10 [=====] - 6s 586ms/step - loss: 0.121
6 - accuracy: 0.9574 - val_loss: 3.1650e-08 - val_accuracy: 1.0000
Epoch 30/40
10/10 [=====] - 5s 543ms/step - loss: 0.104
8 - accuracy: 0.9574 - val_loss: 0.2188 - val_accuracy: 0.9500
Epoch 31/40
10/10 [=====] - 6s 553ms/step - loss: 0.067
5 - accuracy: 0.9468 - val_loss: 0.0084 - val_accuracy: 0.9500
Epoch 32/40
10/10 [=====] - 6s 524ms/step - loss: 0.066
4 - accuracy: 0.9255 - val_loss: 0.2213 - val_accuracy: 0.9500
Epoch 33/40
10/10 [=====] - 5s 540ms/step - loss: 0.114
5 - accuracy: 0.9468 - val_loss: 0.0030 - val_accuracy: 0.9500
Epoch 34/40
10/10 [=====] - 6s 547ms/step - loss: 0.105
6 - accuracy: 0.9468 - val_loss: 0.1134 - val_accuracy: 0.9500
Epoch 35/40
10/10 [=====] - 5s 546ms/step - loss: 0.007
8 - accuracy: 0.9787 - val_loss: 0.1729 - val_accuracy: 0.9000
Epoch 36/40
10/10 [=====] - 5s 583ms/step - loss: 0.249
4 - accuracy: 0.9149 - val_loss: 0.0460 - val_accuracy: 0.9000
Epoch 37/40
10/10 [=====] - 5s 546ms/step - loss: 0.071
8 - accuracy: 0.9362 - val_loss: 0.1263 - val_accuracy: 0.9500
Epoch 38/40
10/10 [=====] - 5s 550ms/step - loss: 0.077
9 - accuracy: 0.9255 - val_loss: 0.0367 - val_accuracy: 0.9500
Epoch 39/40
10/10 [=====] - 5s 548ms/step - loss: 0.043
6 - accuracy: 0.9681 - val_loss: 0.0281 - val_accuracy: 0.9500
Epoch 40/40
10/10 [=====] - 5s 531ms/step - loss: 0.069
0 - accuracy: 0.9149 - val_loss: 1.3706e-06 - val_accuracy: 1.0000
```

[5 points] Plot Accuracy and Loss During Training

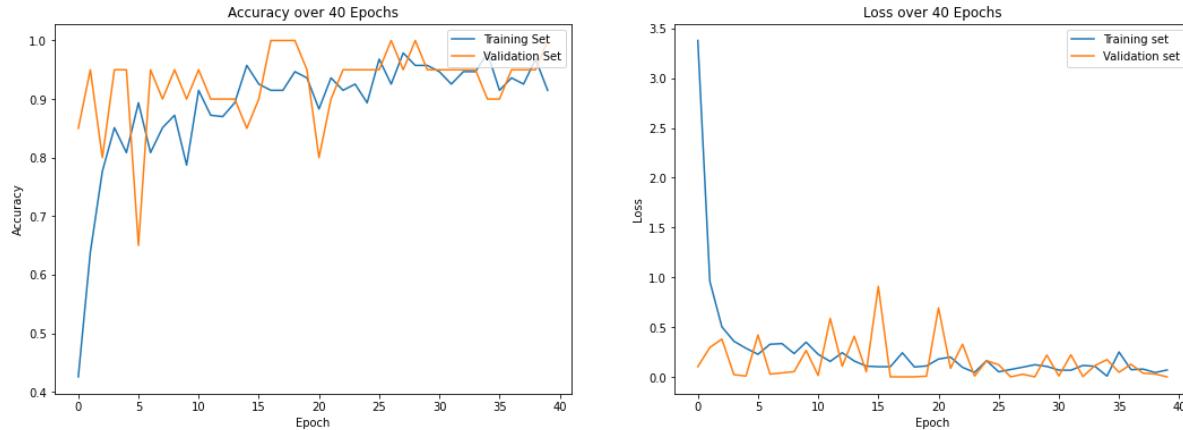
In []:

```
import matplotlib.pyplot as plt
import matplotlib.pyplot as plt
%matplotlib inline

# #Accuracy
plt.figure(figsize=(18,6))
plt.subplot(1,2,1)
plt.plot(history.history["accuracy"])
plt.plot(history.history["val_accuracy"])
plt.title("Accuracy over 40 Epochs")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend(["Training Set", "Validation Set"], loc="upper right")
# Loss
plt.subplot(1,2,2)
plt.plot(history.history["loss"])
plt.plot(history.history["val_loss"])
plt.title("Loss over 40 Epochs")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(["Training set", "Validation set"], loc="upper right")

plt.show()

# raise NotImplementedError("Plot the accuracy and the loss during training")
```

**Plot Test Results**

In []:

```
import matplotlib.image as mpimg

test_datagen = ImageDataGenerator(rescale=1. / 255)
eval_generator = test_datagen.flow_from_directory(TEST_DIR,target_size=IMAGE_SIZE,
                                                 batch_size=1,shuffle=True,seed
=42,class_mode="binary")
eval_generator.reset()
pred = model.predict_generator(eval_generator,18,verbose=1)
for index, probability in enumerate(pred):
    image_path = TEST_DIR + "/" + eval_generator.filenames[index]
    image = mpimg.imread(image_path)
    if image.ndim < 3:
        image = np.reshape(image,(image.shape[0],image.shape[1],1))
        image = np.concatenate([image, image, image], 2)
    #      print(image.shape)

    pixels = np.array(image)
    plt.imshow(pixels)

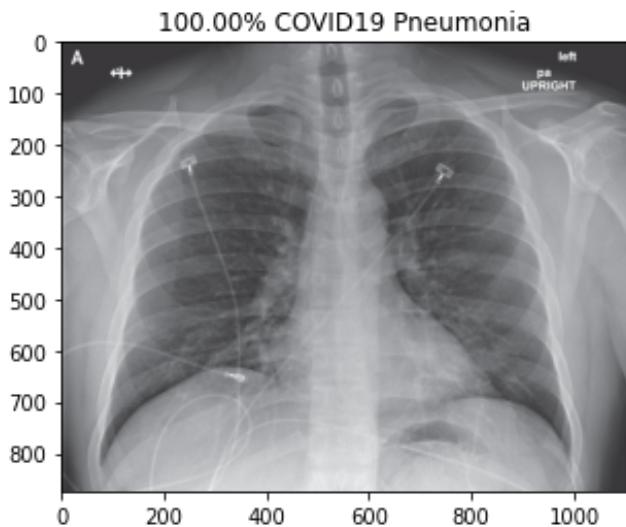
    print(eval_generator.filenames[index])
    if probability > 0.5:
        plt.title("%.2f" % (probability[0]*100) + "% Normal")
    else:
        plt.title("%.2f" % ((1-probability[0])*100) + "% COVID19 Pneumonia")
    plt.show()
```

Found 18 images belonging to 2 classes.

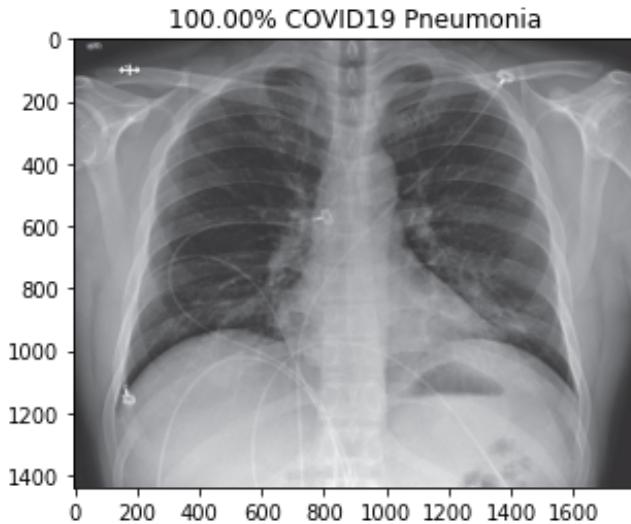
```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:7: User
Warning: `Model.predict_generator` is deprecated and will be removed
in a future version. Please use `Model.predict`, which supports gene
rators.
```

```
import sys
```

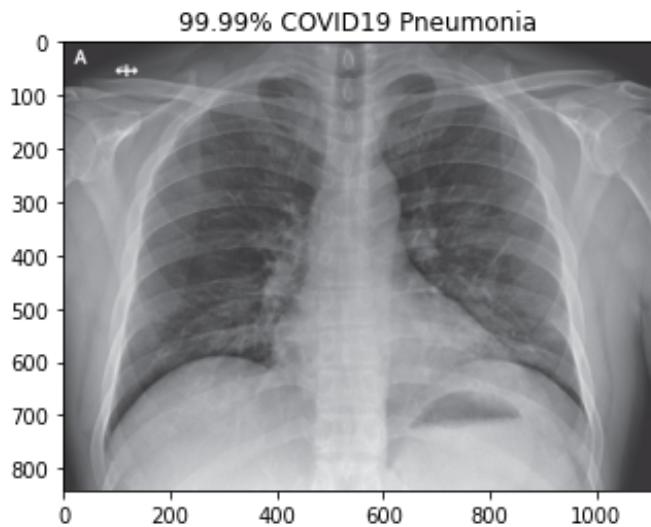
```
18/18 [=====] - 1s 45ms/step
covid/nejmoa2001191_f3-PA.jpeg
```



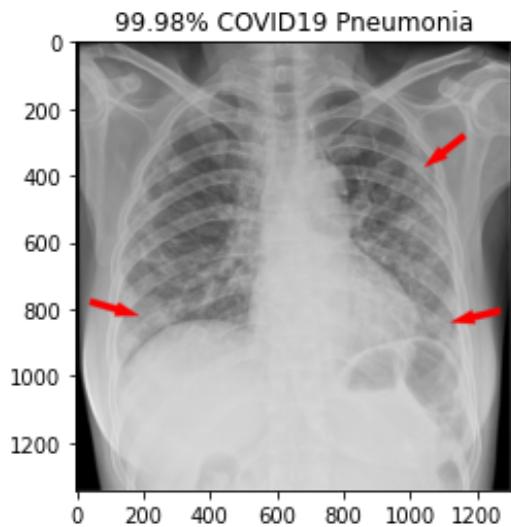
covid/nejmoa2001191_f4.jpeg



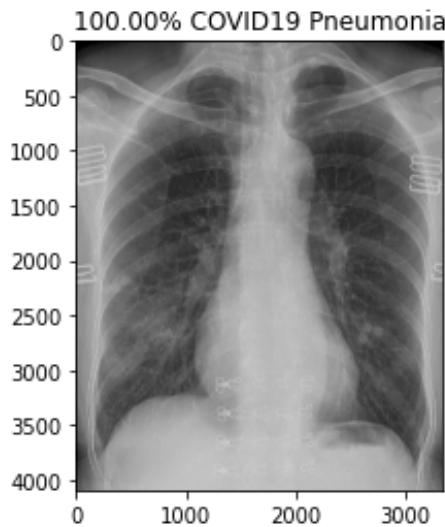
covid/nejmoa2001191_f5-PA.jpeg



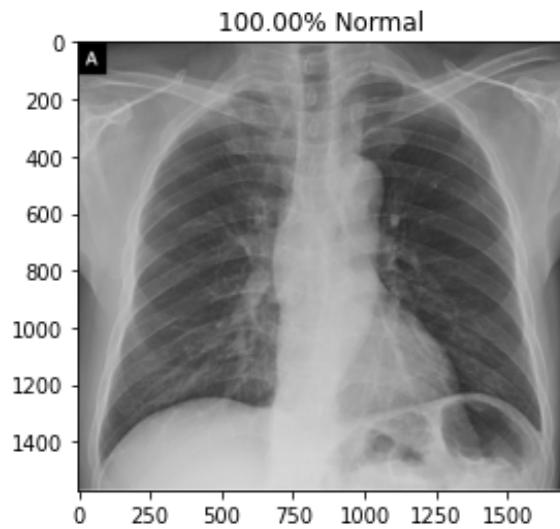
covid/radiol.2020200490.fig3.jpeg



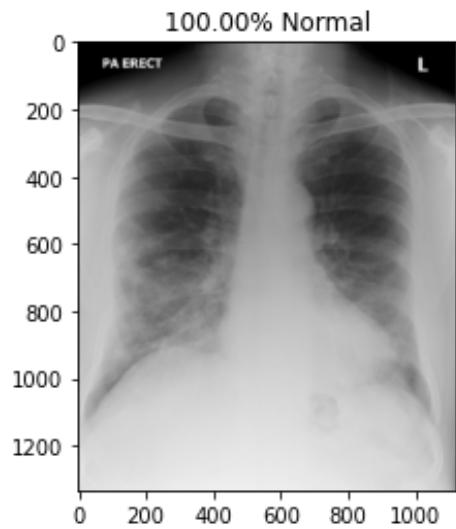
covid/ryct.2020200028.fig1a.jpeg



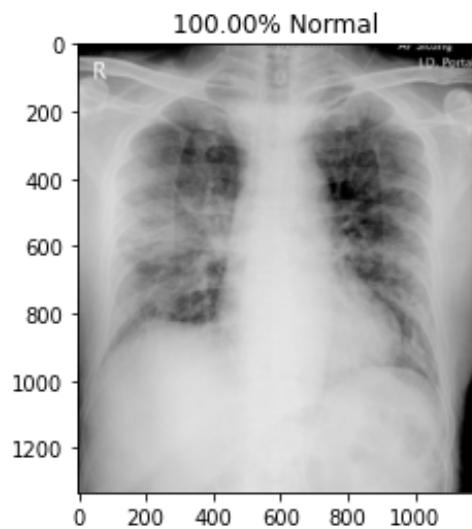
covid/ryct.2020200034.fig2.jpeg



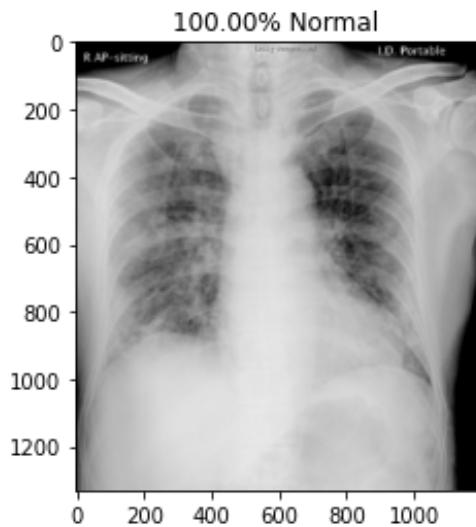
covid/ryct.2020200034.fig5-day0.jpeg



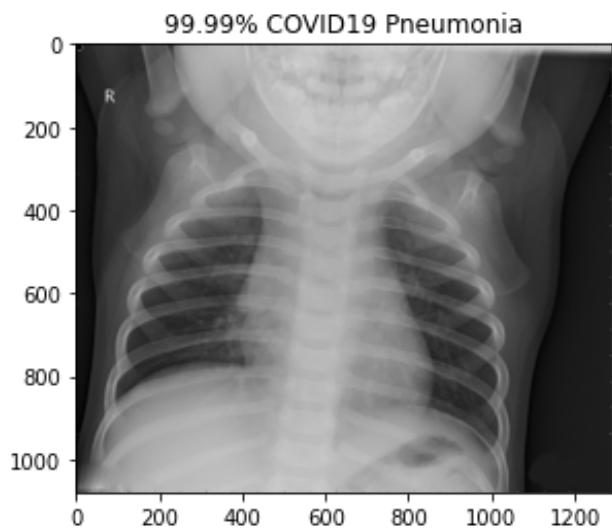
covid/ryct.2020200034.fig5-day4.jpeg



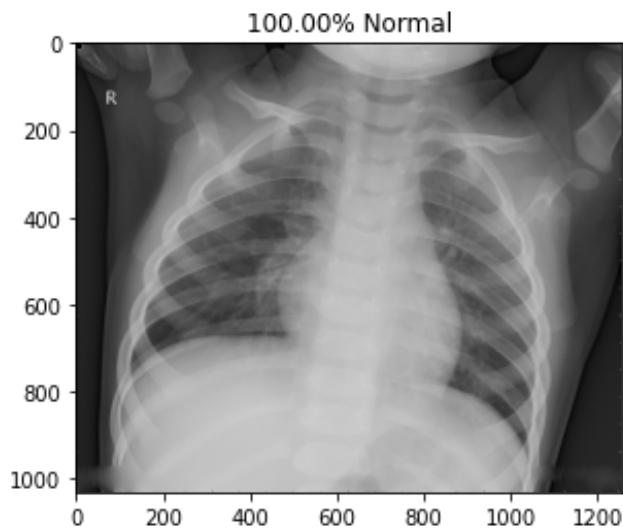
covid/ryct.2020200034.fig5-day7.jpeg



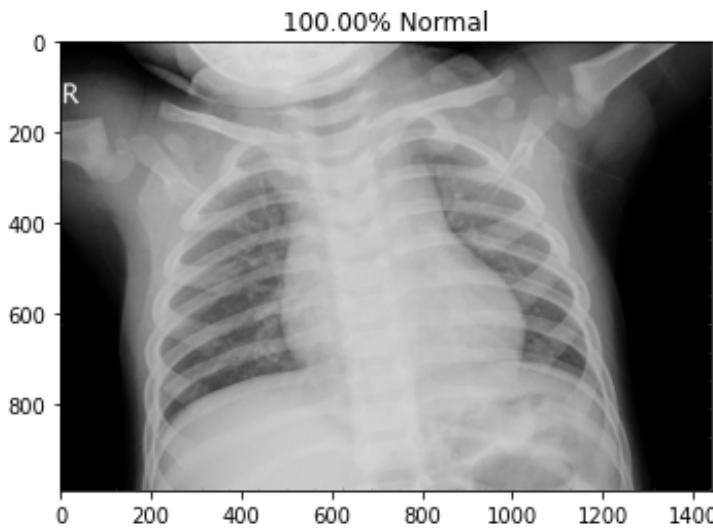
normal/NORMAL2-IM-1385-0001.jpeg



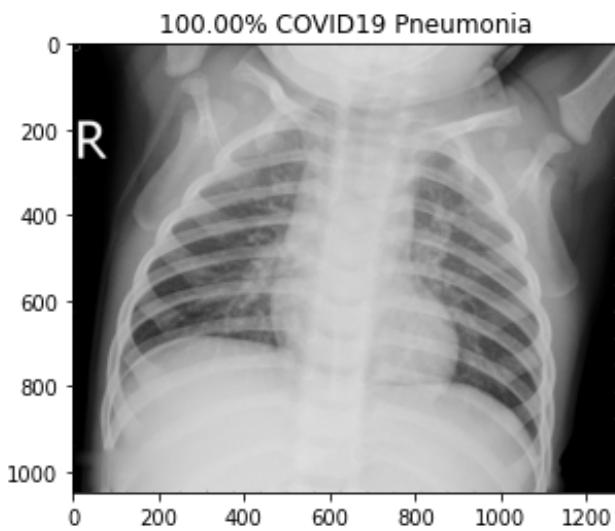
normal/NORMAL2-IM-1396-0001.jpeg



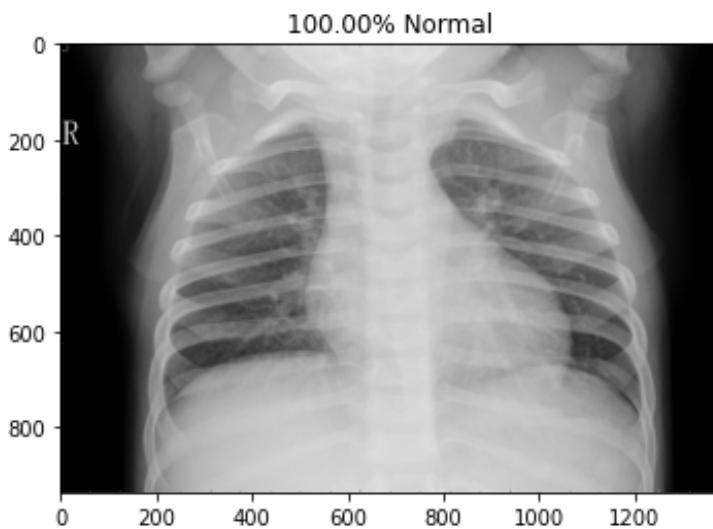
normal/NORMAL2-IM-1400-0001.jpeg



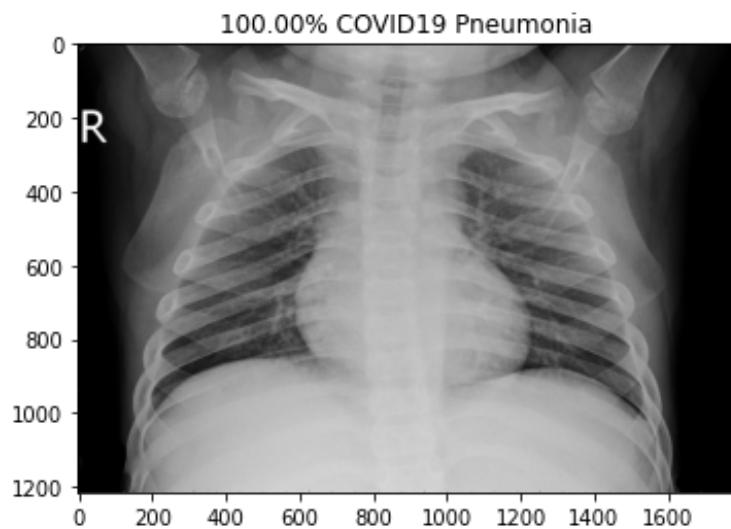
normal/NORMAL2-IM-1401-0001.jpeg



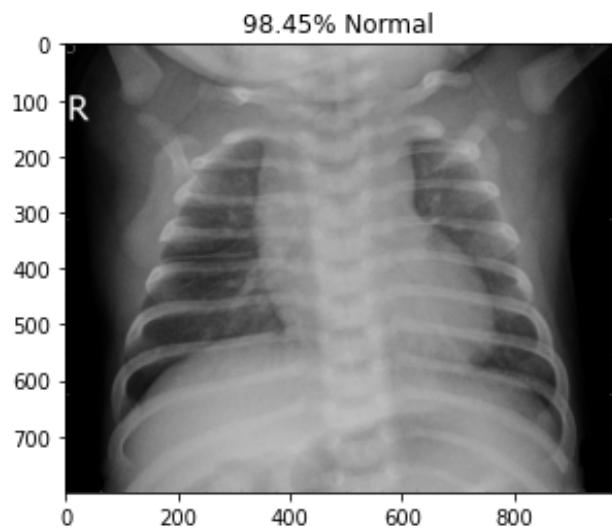
normal/NORMAL2-IM-1406-0001.jpeg



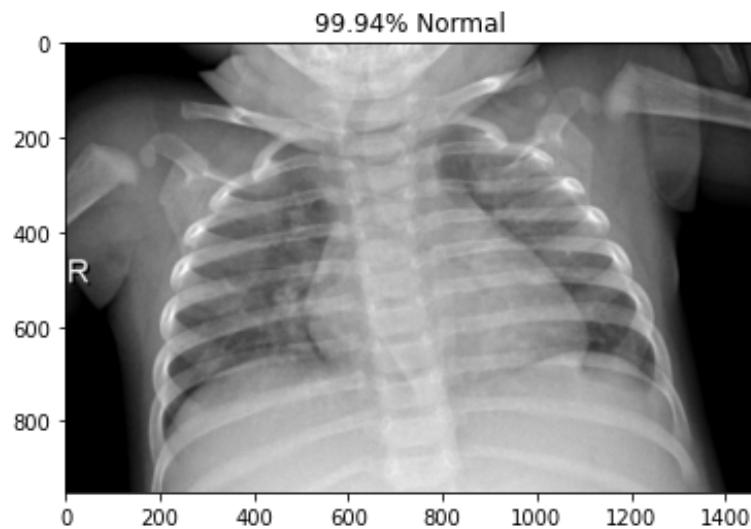
normal/NORMAL2-IM-1412-0001.jpeg



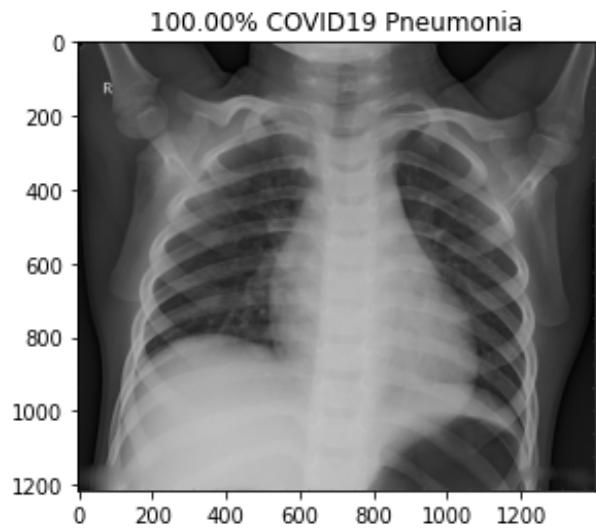
normal/NORMAL2-IM-1419-0001.jpeg



normal/NORMAL2-IM-1422-0001.jpeg



normal/NORMAL2-IM-1423-0001.jpeg



[10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

In []:

```
from sklearn.manifold import TSNE
intermediate_layer_model = models.Model(inputs=model.input,
                                         outputs=model.get_layer('dense_feature')
                                         .output)

tsne_data_generator = test_datagen.flow_from_directory(DATASET_PATH,target_size=
IMAGE_SIZE,
                                                       batch_size=1,shuffle=False,seed
d=42,class_mode="binary")

labels = tsne_data_generator.classes
print(tsne_data_generator.class_indices)

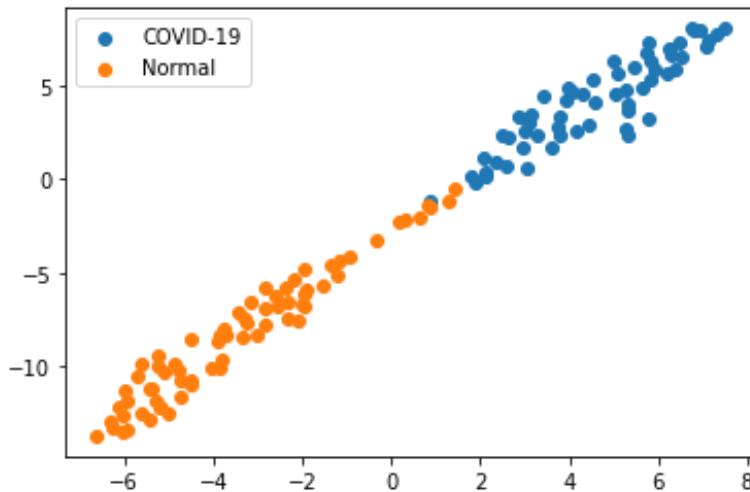
X = TSNE().fit_transform(intermediate_layer_model.predict_generator(tsne_data_ge
nerator, verbose=1))

classes = ["COVID-19", "Normal"]
for i in range(2):
    cluster = X[np.where(labels == i)]
    plt.scatter(cluster[:, 0], cluster[:, 1], label = classes[i])
plt.legend()
```

```
Found 130 images belonging to 2 classes.  
{'covid': 0, 'normal': 1}  
  
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:11: Use  
rWarning: `Model.predict_generator` is deprecated and will be remove  
d in a future version. Please use `Model.predict`, which supports ge  
nerators.  
# This is added back by InteractiveShellApp.init_path()  
  
130/130 [=====] - 4s 29ms/step  
  
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:78  
3: FutureWarning: The default initialization in TSNE will change fro  
m 'random' to 'pca' in 1.2.  
FutureWarning,  
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:79  
3: FutureWarning: The default learning rate in TSNE will change from  
200.0 to 'auto' in 1.2.  
FutureWarning,
```

Out[]:

```
<matplotlib.legend.Legend at 0x7fb2a540c850>
```



In [5]:

```
!jupyter nbconvert --to html '/content/drive/MyDrive/CS 542 Machine Learning/Cha  
llenge/Covid_Data_GradientCrescent/task1_template.ipynb'
```

```
[NbConvertApp] Converting notebook /content/drive/MyDrive/CS 542 Mac  
hine Learning/Challenge/Covid_Data_GradientCrescent/task1_template.i  
pynb to html  
[NbConvertApp] Writing 1685207 bytes to /content/drive/MyDrive/CS 54  
2 Machine Learning/Challenge/Covid_Data_GradientCrescent/task1_temp  
late.html
```

Class Challenge: Image Classification of COVID-19 X-rays

Task 2 [Total points: 30]

In [1]:

```
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))
```

Found GPU at: /device:GPU:0

In [2]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib
conda install numpy
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib
pip install numpy
pip install sklearn
```

Data

Please download the data using the following link: [COVID-19](https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view) (https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view).

- After downloading 'Covid_Data_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
|--all  
|-----train  
|-----test  
|--two  
|-----train  
|-----test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

[20 points] Multi-class Classification

In [55]:

```
import os  
import tensorflow as tf  
import numpy as np  
import matplotlib.pyplot as plt  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from keras.applications.mobilenet import MobileNet  
from keras.applications.vgg16 import VGG16  
from keras import models  
from keras import layers  
from keras import optimizers  
from keras.layers import BatchNormalization, GlobalAveragePooling2D  
from keras.layers.core import Flatten, Dense  
import keras.backend as K  
from sklearn.manifold import TSNE  
os.environ['OMP_NUM_THREADS'] = '1'  
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'  
tf.__version__
```

Out[55]:

'2.8.0'

Load Image Data

In [56]:

```

DATA_LIST = os.listdir('/content/drive/MyDrive/CS 542 Machine Learning/Challenge/Covid_Data_GradientCrescent/all/train')
DATASET_PATH = '/content/drive/MyDrive/CS 542 Machine Learning/Challenge/Covid_Data_GradientCrescent/all/train'
TEST_DIR = '/content/drive/MyDrive/CS 542 Machine Learning/Challenge/Covid_Data_GradientCrescent/all/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU
    runs out of memory
NUM_EPOCHS = 100
LEARNING_RATE = 0.0001 # start off with high rate first 0.001 and experiment with
    reducing it gradually

```

Generate Training and Validation Batches

In [57]:

```

train_datagen = ImageDataGenerator(rescale=1./255, rotation_range=50, featurewise_center = True,
                                    featurewise_std_normalization = True, width_shift_range=0.2,
                                    height_shift_range=0.2, shear_range=0.25, zoom_range=0.1,
                                    zca_whitening = True, channel_shift_range = 20,
                                    horizontal_flip = True, vertical_flip = True,
                                    validation_split = 0.2, fill_mode='constant')

train_batches = train_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                 shuffle=True, batch_size=BATCH_SIZE,
                                                 subset = "training", seed=42,
                                                 class_mode="categorical")

valid_batches = train_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                 shuffle=True, batch_size=BATCH_SIZE,
                                                 subset = "validation",
                                                 seed=42, class_mode="categorical")

/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/image_data_generator.py:342: UserWarning: This ImageDataGenerator specifies `zca_whitening` which overrides setting of `featurewise_std_normalization`.
  warnings.warn('This ImageDataGenerator specifies '
Found 216 images belonging to 4 classes.
Found 54 images belonging to 4 classes.

```

[10 points] Build Model

Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

In [58]:

```
#VGG16
VGG16 = VGG16(weights='imagenet',
                include_top=False,
                input_shape=(224, 224, 3))
VGG16.trainable = False
vgg16 = models.Sequential(name='VGG16')
vgg16.add(VGG16)

vgg16.add(layers.Flatten())
vgg16.add(layers.Dense(256, name='feature_dense', activation='relu'))
vgg16.add(layers.Dense(4, activation='softmax'))
vgg16.summary()

vgg16.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0005), loss='categorical_crossentropy', metrics=['acc'])
```

Model: "VGG16"

Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_12 (Flatten)	(None, 25088)	0
feature_dense (Dense)	(None, 256)	6422784
dense_11 (Dense)	(None, 4)	1028
<hr/>		
Total params: 21,138,500		
Trainable params: 6,423,812		
Non-trainable params: 14,714,688		

In [48]:

```
# MobileNet

mobile = MobileNet(include_top=False, weights='imagenet', input_shape=(224, 224, 3))

mobilenet = models.Sequential(name='MobileNet')
mobilenet.add(mobile)
mobilenet.add(GlobalAveragePooling2D())
mobilenet.add(layers.Flatten())
mobilenet.add(layers.Dense(512, activation='relu'))
mobilenet.add(layers.Dense(256, name='feature_dense'))

mobilenet.add(layers.Dense(4, activation='softmax'))

mobilenet.summary()
mobilenet.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001), loss='categorical_crossentropy', metrics=[ "acc" ])
```

Model: "MobileNet"

Layer (type)	Output Shape	Param #
<hr/>		
mobilenet_1.00_224 (Functional)	(None, 7, 7, 1024)	3228864
global_average_pooling2d_5 (GlobalAveragePooling2D)	(None, 1024)	0
flatten_11 (Flatten)	(None, 1024)	0
dense_9 (Dense)	(None, 512)	524800
feature_dense (Dense)	(None, 256)	131328
dense_10 (Dense)	(None, 4)	1028
<hr/>		
Total params: 3,886,020		
Trainable params: 3,864,132		
Non-trainable params: 21,888		

[5 points] Train Model

In [59]:

```
#FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN=train_batches.n//train_batches.batch_size
STEP_SIZE_VALID=valid_batches.n//valid_batches.batch_size
```

22
6

In [60]:

```
#VGG16
vgg16_model = vgg16.fit(train_batches,
                        steps_per_epoch=STEP_SIZE_TRAIN,
                        epochs = 100,
                        validation_data=valid_batches,
                        validation_steps=STEP_SIZE_VALID
)
# shuffle=True
# raise NotImplementedError("Use the model.fit function to train your network")
```

Epoch 1/100
21/21 [=====] - 9s 411ms/step - loss: 2.159
4 - acc: 0.3301 - val_loss: 1.2102 - val_acc: 0.4800
Epoch 2/100
21/21 [=====] - 8s 390ms/step - loss: 1.213
7 - acc: 0.4757 - val_loss: 0.9627 - val_acc: 0.6400
Epoch 3/100
21/21 [=====] - 8s 384ms/step - loss: 1.088
2 - acc: 0.4951 - val_loss: 0.8244 - val_acc: 0.7200
Epoch 4/100
21/21 [=====] - 8s 390ms/step - loss: 1.027
6 - acc: 0.5728 - val_loss: 0.9864 - val_acc: 0.5200
Epoch 5/100
21/21 [=====] - 8s 389ms/step - loss: 0.988
2 - acc: 0.6165 - val_loss: 0.8828 - val_acc: 0.6600
Epoch 6/100
21/21 [=====] - 8s 387ms/step - loss: 1.075
4 - acc: 0.5680 - val_loss: 0.9325 - val_acc: 0.5200
Epoch 7/100
21/21 [=====] - 8s 390ms/step - loss: 0.905
3 - acc: 0.5680 - val_loss: 0.8394 - val_acc: 0.5800
Epoch 8/100
21/21 [=====] - 8s 377ms/step - loss: 0.867
2 - acc: 0.6311 - val_loss: 0.8822 - val_acc: 0.6200
Epoch 9/100
21/21 [=====] - 8s 395ms/step - loss: 0.901
2 - acc: 0.6505 - val_loss: 0.7763 - val_acc: 0.6200
Epoch 10/100
21/21 [=====] - 8s 387ms/step - loss: 0.871
1 - acc: 0.6505 - val_loss: 0.7291 - val_acc: 0.6200
Epoch 11/100
21/21 [=====] - 8s 380ms/step - loss: 0.826
1 - acc: 0.6650 - val_loss: 0.8383 - val_acc: 0.6200
Epoch 12/100
21/21 [=====] - 8s 381ms/step - loss: 0.782
2 - acc: 0.6942 - val_loss: 0.9880 - val_acc: 0.5600
Epoch 13/100
21/21 [=====] - 8s 395ms/step - loss: 0.795
9 - acc: 0.6505 - val_loss: 0.7630 - val_acc: 0.6400
Epoch 14/100
21/21 [=====] - 8s 388ms/step - loss: 0.795
7 - acc: 0.6602 - val_loss: 0.6895 - val_acc: 0.6600
Epoch 15/100
21/21 [=====] - 8s 393ms/step - loss: 0.677
6 - acc: 0.7039 - val_loss: 0.6999 - val_acc: 0.7200
Epoch 16/100
21/21 [=====] - 8s 391ms/step - loss: 0.752
8 - acc: 0.7039 - val_loss: 0.7791 - val_acc: 0.6400
Epoch 17/100
21/21 [=====] - 8s 394ms/step - loss: 0.744
9 - acc: 0.6553 - val_loss: 0.7346 - val_acc: 0.7000
Epoch 18/100
21/21 [=====] - 8s 387ms/step - loss: 0.744
4 - acc: 0.6699 - val_loss: 0.9277 - val_acc: 0.5200
Epoch 19/100
21/21 [=====] - 8s 385ms/step - loss: 0.781
7 - acc: 0.7087 - val_loss: 0.7715 - val_acc: 0.7000
Epoch 20/100
21/21 [=====] - 8s 391ms/step - loss: 0.679
5 - acc: 0.7087 - val_loss: 0.9009 - val_acc: 0.6200
Epoch 21/100

```
21/21 [=====] - 8s 387ms/step - loss: 0.796
5 - acc: 0.6602 - val_loss: 0.7138 - val_acc: 0.6600
Epoch 22/100
21/21 [=====] - 8s 389ms/step - loss: 0.714
3 - acc: 0.7233 - val_loss: 0.8226 - val_acc: 0.5800
Epoch 23/100
21/21 [=====] - 8s 384ms/step - loss: 0.694
2 - acc: 0.6942 - val_loss: 0.8100 - val_acc: 0.6400
Epoch 24/100
21/21 [=====] - 8s 391ms/step - loss: 0.708
8 - acc: 0.7184 - val_loss: 0.7186 - val_acc: 0.6400
Epoch 25/100
21/21 [=====] - 8s 387ms/step - loss: 0.629
4 - acc: 0.7184 - val_loss: 0.7841 - val_acc: 0.6800
Epoch 26/100
21/21 [=====] - 8s 394ms/step - loss: 0.693
6 - acc: 0.6845 - val_loss: 0.7016 - val_acc: 0.7200
Epoch 27/100
21/21 [=====] - 8s 400ms/step - loss: 0.669
6 - acc: 0.6942 - val_loss: 0.6036 - val_acc: 0.7200
Epoch 28/100
21/21 [=====] - 8s 393ms/step - loss: 0.673
2 - acc: 0.7427 - val_loss: 0.7567 - val_acc: 0.6000
Epoch 29/100
21/21 [=====] - 8s 388ms/step - loss: 0.626
0 - acc: 0.7670 - val_loss: 0.7274 - val_acc: 0.6600
Epoch 30/100
21/21 [=====] - 8s 392ms/step - loss: 0.705
4 - acc: 0.6893 - val_loss: 0.9401 - val_acc: 0.5600
Epoch 31/100
21/21 [=====] - 8s 383ms/step - loss: 0.732
4 - acc: 0.6845 - val_loss: 0.5446 - val_acc: 0.7000
Epoch 32/100
21/21 [=====] - 8s 390ms/step - loss: 0.592
9 - acc: 0.7427 - val_loss: 0.7379 - val_acc: 0.6800
Epoch 33/100
21/21 [=====] - 8s 392ms/step - loss: 0.635
0 - acc: 0.7524 - val_loss: 0.9842 - val_acc: 0.5200
Epoch 34/100
21/21 [=====] - 8s 388ms/step - loss: 0.582
8 - acc: 0.7476 - val_loss: 0.6998 - val_acc: 0.6600
Epoch 35/100
21/21 [=====] - 8s 396ms/step - loss: 0.583
0 - acc: 0.7670 - val_loss: 0.8025 - val_acc: 0.6400
Epoch 36/100
21/21 [=====] - 8s 392ms/step - loss: 0.714
9 - acc: 0.7190 - val_loss: 0.8865 - val_acc: 0.6600
Epoch 37/100
21/21 [=====] - 8s 390ms/step - loss: 0.650
2 - acc: 0.6796 - val_loss: 0.7115 - val_acc: 0.6600
Epoch 38/100
21/21 [=====] - 8s 395ms/step - loss: 0.631
3 - acc: 0.7184 - val_loss: 0.9736 - val_acc: 0.6400
Epoch 39/100
21/21 [=====] - 8s 389ms/step - loss: 0.583
3 - acc: 0.7476 - val_loss: 0.7085 - val_acc: 0.6600
Epoch 40/100
21/21 [=====] - 8s 393ms/step - loss: 0.626
0 - acc: 0.7379 - val_loss: 0.7374 - val_acc: 0.7000
Epoch 41/100
21/21 [=====] - 9s 444ms/step - loss: 0.552
```

```
5 - acc: 0.7718 - val_loss: 0.7911 - val_acc: 0.6400
Epoch 42/100
21/21 [=====] - 8s 385ms/step - loss: 0.660
2 - acc: 0.7233 - val_loss: 0.6949 - val_acc: 0.7000
Epoch 43/100
21/21 [=====] - 8s 390ms/step - loss: 0.695
8 - acc: 0.7184 - val_loss: 0.6801 - val_acc: 0.6400
Epoch 44/100
21/21 [=====] - 8s 392ms/step - loss: 0.593
7 - acc: 0.7330 - val_loss: 0.5773 - val_acc: 0.6800
Epoch 45/100
21/21 [=====] - 8s 391ms/step - loss: 0.587
9 - acc: 0.7864 - val_loss: 0.9509 - val_acc: 0.5400
Epoch 46/100
21/21 [=====] - 8s 392ms/step - loss: 0.591
0 - acc: 0.7330 - val_loss: 0.9701 - val_acc: 0.5600
Epoch 47/100
21/21 [=====] - 8s 383ms/step - loss: 0.587
6 - acc: 0.7476 - val_loss: 0.6868 - val_acc: 0.7400
Epoch 48/100
21/21 [=====] - 8s 397ms/step - loss: 0.626
0 - acc: 0.7379 - val_loss: 0.8409 - val_acc: 0.6000
Epoch 49/100
21/21 [=====] - 8s 392ms/step - loss: 0.599
6 - acc: 0.7233 - val_loss: 0.8180 - val_acc: 0.6400
Epoch 50/100
21/21 [=====] - 8s 395ms/step - loss: 0.596
4 - acc: 0.7524 - val_loss: 0.7686 - val_acc: 0.6400
Epoch 51/100
21/21 [=====] - 8s 399ms/step - loss: 0.584
3 - acc: 0.7767 - val_loss: 0.7334 - val_acc: 0.7000
Epoch 52/100
21/21 [=====] - 8s 388ms/step - loss: 0.602
4 - acc: 0.7524 - val_loss: 0.6984 - val_acc: 0.7200
Epoch 53/100
21/21 [=====] - 8s 388ms/step - loss: 0.554
6 - acc: 0.7718 - val_loss: 0.6975 - val_acc: 0.7400
Epoch 54/100
21/21 [=====] - 8s 391ms/step - loss: 0.532
2 - acc: 0.7718 - val_loss: 0.7302 - val_acc: 0.7000
Epoch 55/100
21/21 [=====] - 8s 388ms/step - loss: 0.508
7 - acc: 0.8107 - val_loss: 0.6479 - val_acc: 0.6800
Epoch 56/100
21/21 [=====] - 8s 392ms/step - loss: 0.553
0 - acc: 0.7718 - val_loss: 0.6737 - val_acc: 0.6200
Epoch 57/100
21/21 [=====] - 8s 390ms/step - loss: 0.565
3 - acc: 0.7670 - val_loss: 0.7575 - val_acc: 0.6200
Epoch 58/100
21/21 [=====] - 8s 386ms/step - loss: 0.581
9 - acc: 0.7670 - val_loss: 0.6804 - val_acc: 0.7200
Epoch 59/100
21/21 [=====] - 8s 388ms/step - loss: 0.551
5 - acc: 0.7379 - val_loss: 0.7621 - val_acc: 0.7000
Epoch 60/100
21/21 [=====] - 8s 384ms/step - loss: 0.558
7 - acc: 0.7718 - val_loss: 0.7091 - val_acc: 0.6000
Epoch 61/100
21/21 [=====] - 8s 393ms/step - loss: 0.627
4 - acc: 0.7282 - val_loss: 0.7467 - val_acc: 0.6800
```

Epoch 62/100
21/21 [=====] - 8s 391ms/step - loss: 0.517
3 - acc: 0.7767 - val_loss: 0.5992 - val_acc: 0.6800
Epoch 63/100
21/21 [=====] - 8s 396ms/step - loss: 0.560
4 - acc: 0.7767 - val_loss: 0.6629 - val_acc: 0.6600
Epoch 64/100
21/21 [=====] - 8s 389ms/step - loss: 0.501
4 - acc: 0.7913 - val_loss: 0.7208 - val_acc: 0.6600
Epoch 65/100
21/21 [=====] - 8s 396ms/step - loss: 0.528
3 - acc: 0.7476 - val_loss: 0.7954 - val_acc: 0.6800
Epoch 66/100
21/21 [=====] - 8s 393ms/step - loss: 0.508
9 - acc: 0.8048 - val_loss: 0.7439 - val_acc: 0.7200
Epoch 67/100
21/21 [=====] - 8s 393ms/step - loss: 0.647
7 - acc: 0.7282 - val_loss: 0.7253 - val_acc: 0.6400
Epoch 68/100
21/21 [=====] - 8s 393ms/step - loss: 0.548
4 - acc: 0.7573 - val_loss: 0.8718 - val_acc: 0.4800
Epoch 69/100
21/21 [=====] - 8s 390ms/step - loss: 0.541
9 - acc: 0.7670 - val_loss: 0.8435 - val_acc: 0.5400
Epoch 70/100
21/21 [=====] - 8s 394ms/step - loss: 0.471
4 - acc: 0.8252 - val_loss: 0.6871 - val_acc: 0.6600
Epoch 71/100
21/21 [=====] - 8s 388ms/step - loss: 0.603
9 - acc: 0.7476 - val_loss: 0.5000 - val_acc: 0.7400
Epoch 72/100
21/21 [=====] - 8s 395ms/step - loss: 0.587
4 - acc: 0.7767 - val_loss: 0.6767 - val_acc: 0.7000
Epoch 73/100
21/21 [=====] - 8s 389ms/step - loss: 0.614
5 - acc: 0.7330 - val_loss: 0.9293 - val_acc: 0.5400
Epoch 74/100
21/21 [=====] - 8s 386ms/step - loss: 0.537
9 - acc: 0.7427 - val_loss: 0.6488 - val_acc: 0.7000
Epoch 75/100
21/21 [=====] - 8s 393ms/step - loss: 0.551
2 - acc: 0.7961 - val_loss: 0.8361 - val_acc: 0.6600
Epoch 76/100
21/21 [=====] - 8s 395ms/step - loss: 0.533
6 - acc: 0.7816 - val_loss: 0.5689 - val_acc: 0.6800
Epoch 77/100
21/21 [=====] - 8s 394ms/step - loss: 0.446
5 - acc: 0.8252 - val_loss: 0.5762 - val_acc: 0.7200
Epoch 78/100
21/21 [=====] - 8s 382ms/step - loss: 0.546
5 - acc: 0.7718 - val_loss: 0.6611 - val_acc: 0.6600
Epoch 79/100
21/21 [=====] - 8s 385ms/step - loss: 0.523
9 - acc: 0.7524 - val_loss: 0.6583 - val_acc: 0.7200
Epoch 80/100
21/21 [=====] - 8s 391ms/step - loss: 0.588
6 - acc: 0.7379 - val_loss: 0.8412 - val_acc: 0.6600
Epoch 81/100
21/21 [=====] - 8s 390ms/step - loss: 0.475
0 - acc: 0.8155 - val_loss: 0.8145 - val_acc: 0.6600
Epoch 82/100

```
21/21 [=====] - 8s 386ms/step - loss: 0.474
5 - acc: 0.7961 - val_loss: 0.6249 - val_acc: 0.6600
Epoch 83/100
21/21 [=====] - 8s 395ms/step - loss: 0.463
7 - acc: 0.8010 - val_loss: 0.8369 - val_acc: 0.6400
Epoch 84/100
21/21 [=====] - 8s 385ms/step - loss: 0.520
1 - acc: 0.7718 - val_loss: 0.6279 - val_acc: 0.6400
Epoch 85/100
21/21 [=====] - 8s 392ms/step - loss: 0.473
9 - acc: 0.7952 - val_loss: 0.7062 - val_acc: 0.6000
Epoch 86/100
21/21 [=====] - 8s 393ms/step - loss: 0.502
7 - acc: 0.7816 - val_loss: 1.0830 - val_acc: 0.5400
Epoch 87/100
21/21 [=====] - 8s 394ms/step - loss: 0.545
1 - acc: 0.7379 - val_loss: 0.9031 - val_acc: 0.6200
Epoch 88/100
21/21 [=====] - 8s 386ms/step - loss: 0.450
9 - acc: 0.7864 - val_loss: 0.5349 - val_acc: 0.7000
Epoch 89/100
21/21 [=====] - 8s 392ms/step - loss: 0.574
5 - acc: 0.7857 - val_loss: 0.5987 - val_acc: 0.6800
Epoch 90/100
21/21 [=====] - 8s 395ms/step - loss: 0.599
6 - acc: 0.7621 - val_loss: 0.6162 - val_acc: 0.6800
Epoch 91/100
21/21 [=====] - 8s 388ms/step - loss: 0.511
8 - acc: 0.7524 - val_loss: 0.6943 - val_acc: 0.6800
Epoch 92/100
21/21 [=====] - 8s 376ms/step - loss: 0.577
9 - acc: 0.7282 - val_loss: 0.7491 - val_acc: 0.6600
Epoch 93/100
21/21 [=====] - 8s 381ms/step - loss: 0.478
2 - acc: 0.7816 - val_loss: 0.6965 - val_acc: 0.6800
Epoch 94/100
21/21 [=====] - 8s 387ms/step - loss: 0.484
8 - acc: 0.8107 - val_loss: 0.6054 - val_acc: 0.6800
Epoch 95/100
21/21 [=====] - 8s 389ms/step - loss: 0.572
7 - acc: 0.7476 - val_loss: 1.0214 - val_acc: 0.6000
Epoch 96/100
21/21 [=====] - 9s 421ms/step - loss: 0.574
6 - acc: 0.7670 - val_loss: 0.6953 - val_acc: 0.6600
Epoch 97/100
21/21 [=====] - 8s 383ms/step - loss: 0.462
0 - acc: 0.8398 - val_loss: 0.5763 - val_acc: 0.7000
Epoch 98/100
21/21 [=====] - 8s 395ms/step - loss: 0.475
7 - acc: 0.8010 - val_loss: 0.6667 - val_acc: 0.7000
Epoch 99/100
21/21 [=====] - 8s 400ms/step - loss: 0.508
5 - acc: 0.8010 - val_loss: 0.6881 - val_acc: 0.6600
Epoch 100/100
21/21 [=====] - 8s 398ms/step - loss: 0.481
7 - acc: 0.7864 - val_loss: 0.5624 - val_acc: 0.7200
```

In [26]:

```
#MobileNet
mobilenet_model=mobilenet.fit(train_batches,
                               steps_per_epoch=STEP_SIZE_TRAIN,
                               epochs = 100,
                               validation_data=valid_batches,
                               validation_steps=STEP_SIZE_VALID,
                               shuffle=True)
```

Epoch 1/100
21/21 [=====] - 13s 451ms/step - loss: 1.21
23 - acc: 0.4417 - val_loss: 2.4747 - val_acc: 0.2600
Epoch 2/100
21/21 [=====] - 8s 401ms/step - loss: 0.705
6 - acc: 0.7136 - val_loss: 2.9546 - val_acc: 0.2000
Epoch 3/100
21/21 [=====] - 9s 402ms/step - loss: 0.615
9 - acc: 0.7282 - val_loss: 2.9787 - val_acc: 0.3000
Epoch 4/100
21/21 [=====] - 9s 405ms/step - loss: 0.561
6 - acc: 0.7913 - val_loss: 3.0249 - val_acc: 0.2400
Epoch 5/100
21/21 [=====] - 8s 403ms/step - loss: 0.549
6 - acc: 0.7767 - val_loss: 2.3122 - val_acc: 0.3000
Epoch 6/100
21/21 [=====] - 8s 401ms/step - loss: 0.531
3 - acc: 0.7621 - val_loss: 2.3439 - val_acc: 0.3000
Epoch 7/100
21/21 [=====] - 8s 399ms/step - loss: 0.442
2 - acc: 0.8204 - val_loss: 1.6809 - val_acc: 0.4200
Epoch 8/100
21/21 [=====] - 8s 399ms/step - loss: 0.494
4 - acc: 0.7816 - val_loss: 2.5028 - val_acc: 0.3600
Epoch 9/100
21/21 [=====] - 8s 402ms/step - loss: 0.438
0 - acc: 0.8350 - val_loss: 1.7472 - val_acc: 0.3600
Epoch 10/100
21/21 [=====] - 10s 460ms/step - loss: 0.37
02 - acc: 0.8738 - val_loss: 2.2474 - val_acc: 0.4200
Epoch 11/100
21/21 [=====] - 8s 406ms/step - loss: 0.392
1 - acc: 0.8495 - val_loss: 2.3004 - val_acc: 0.3600
Epoch 12/100
21/21 [=====] - 8s 399ms/step - loss: 0.415
6 - acc: 0.8204 - val_loss: 1.5798 - val_acc: 0.4400
Epoch 13/100
21/21 [=====] - 8s 394ms/step - loss: 0.374
7 - acc: 0.8495 - val_loss: 1.2463 - val_acc: 0.5200
Epoch 14/100
21/21 [=====] - 8s 405ms/step - loss: 0.385
4 - acc: 0.8398 - val_loss: 1.6831 - val_acc: 0.4400
Epoch 15/100
21/21 [=====] - 8s 400ms/step - loss: 0.454
0 - acc: 0.8252 - val_loss: 2.2206 - val_acc: 0.3200
Epoch 16/100
21/21 [=====] - 8s 401ms/step - loss: 0.294
2 - acc: 0.8786 - val_loss: 2.0693 - val_acc: 0.3600
Epoch 17/100
21/21 [=====] - 8s 407ms/step - loss: 0.406
0 - acc: 0.8301 - val_loss: 1.6325 - val_acc: 0.4400
Epoch 18/100
21/21 [=====] - 8s 404ms/step - loss: 0.306
2 - acc: 0.8592 - val_loss: 1.5158 - val_acc: 0.4200
Epoch 19/100
21/21 [=====] - 8s 399ms/step - loss: 0.285
2 - acc: 0.8738 - val_loss: 0.8467 - val_acc: 0.5600
Epoch 20/100
21/21 [=====] - 8s 396ms/step - loss: 0.284
9 - acc: 0.8689 - val_loss: 1.1628 - val_acc: 0.5600
Epoch 21/100

```
21/21 [=====] - 8s 394ms/step - loss: 0.235
8 - acc: 0.9078 - val_loss: 1.6709 - val_acc: 0.4200
Epoch 22/100
21/21 [=====] - 8s 403ms/step - loss: 0.259
1 - acc: 0.9029 - val_loss: 1.6148 - val_acc: 0.5000
Epoch 23/100
21/21 [=====] - 8s 402ms/step - loss: 0.273
4 - acc: 0.9029 - val_loss: 1.2389 - val_acc: 0.6000
Epoch 24/100
21/21 [=====] - 8s 395ms/step - loss: 0.327
6 - acc: 0.8738 - val_loss: 0.9777 - val_acc: 0.6200
Epoch 25/100
21/21 [=====] - 8s 395ms/step - loss: 0.251
2 - acc: 0.9126 - val_loss: 0.9395 - val_acc: 0.6600
Epoch 26/100
21/21 [=====] - 8s 404ms/step - loss: 0.266
3 - acc: 0.8883 - val_loss: 0.6061 - val_acc: 0.7200
Epoch 27/100
21/21 [=====] - 8s 404ms/step - loss: 0.248
6 - acc: 0.8932 - val_loss: 0.6149 - val_acc: 0.7800
Epoch 28/100
21/21 [=====] - 8s 403ms/step - loss: 0.164
2 - acc: 0.9320 - val_loss: 0.4695 - val_acc: 0.7800
Epoch 29/100
21/21 [=====] - 8s 403ms/step - loss: 0.164
0 - acc: 0.9466 - val_loss: 0.5033 - val_acc: 0.8200
Epoch 30/100
21/21 [=====] - 8s 395ms/step - loss: 0.254
6 - acc: 0.9126 - val_loss: 0.7416 - val_acc: 0.6600
Epoch 31/100
21/21 [=====] - 8s 391ms/step - loss: 0.302
0 - acc: 0.8592 - val_loss: 0.5365 - val_acc: 0.7600
Epoch 32/100
21/21 [=====] - 8s 397ms/step - loss: 0.169
2 - acc: 0.9078 - val_loss: 0.9456 - val_acc: 0.6800
Epoch 33/100
21/21 [=====] - 8s 402ms/step - loss: 0.175
5 - acc: 0.9369 - val_loss: 0.7911 - val_acc: 0.7200
Epoch 34/100
21/21 [=====] - 8s 391ms/step - loss: 0.080
5 - acc: 0.9854 - val_loss: 0.5854 - val_acc: 0.7800
Epoch 35/100
21/21 [=====] - 8s 397ms/step - loss: 0.154
5 - acc: 0.9417 - val_loss: 1.0097 - val_acc: 0.7800
Epoch 36/100
21/21 [=====] - 9s 413ms/step - loss: 0.202
1 - acc: 0.9143 - val_loss: 0.8347 - val_acc: 0.7800
Epoch 37/100
21/21 [=====] - 8s 398ms/step - loss: 0.170
7 - acc: 0.9272 - val_loss: 0.8581 - val_acc: 0.6600
Epoch 38/100
21/21 [=====] - 8s 398ms/step - loss: 0.156
2 - acc: 0.9320 - val_loss: 0.5998 - val_acc: 0.8000
Epoch 39/100
21/21 [=====] - 8s 408ms/step - loss: 0.234
0 - acc: 0.9320 - val_loss: 0.7261 - val_acc: 0.7200
Epoch 40/100
21/21 [=====] - 8s 399ms/step - loss: 0.127
1 - acc: 0.9417 - val_loss: 0.6601 - val_acc: 0.7600
Epoch 41/100
21/21 [=====] - 8s 404ms/step - loss: 0.111
```

```
3 - acc: 0.9714 - val_loss: 0.4811 - val_acc: 0.8000
Epoch 42/100
21/21 [=====] - 8s 397ms/step - loss: 0.068
2 - acc: 0.9757 - val_loss: 0.7911 - val_acc: 0.7400
Epoch 43/100
21/21 [=====] - 8s 398ms/step - loss: 0.089
3 - acc: 0.9709 - val_loss: 0.6227 - val_acc: 0.8000
Epoch 44/100
21/21 [=====] - 8s 393ms/step - loss: 0.109
1 - acc: 0.9476 - val_loss: 1.1026 - val_acc: 0.7200
Epoch 45/100
21/21 [=====] - 8s 400ms/step - loss: 0.182
2 - acc: 0.9320 - val_loss: 1.0055 - val_acc: 0.6600
Epoch 46/100
21/21 [=====] - 8s 396ms/step - loss: 0.166
6 - acc: 0.9320 - val_loss: 0.7135 - val_acc: 0.8000
Epoch 47/100
21/21 [=====] - 8s 394ms/step - loss: 0.219
7 - acc: 0.9272 - val_loss: 0.7175 - val_acc: 0.7800
Epoch 48/100
21/21 [=====] - 8s 401ms/step - loss: 0.118
9 - acc: 0.9563 - val_loss: 0.6342 - val_acc: 0.8000
Epoch 49/100
21/21 [=====] - 8s 405ms/step - loss: 0.112
6 - acc: 0.9612 - val_loss: 0.7755 - val_acc: 0.7800
Epoch 50/100
21/21 [=====] - 8s 408ms/step - loss: 0.166
5 - acc: 0.9272 - val_loss: 0.6273 - val_acc: 0.7800
Epoch 51/100
21/21 [=====] - 8s 404ms/step - loss: 0.228
7 - acc: 0.9095 - val_loss: 0.4424 - val_acc: 0.8400
Epoch 52/100
21/21 [=====] - 8s 398ms/step - loss: 0.162
6 - acc: 0.9320 - val_loss: 0.8378 - val_acc: 0.7400
Epoch 53/100
21/21 [=====] - 8s 391ms/step - loss: 0.101
7 - acc: 0.9660 - val_loss: 0.7555 - val_acc: 0.8000
Epoch 54/100
21/21 [=====] - 8s 405ms/step - loss: 0.125
2 - acc: 0.9515 - val_loss: 1.1830 - val_acc: 0.7000
Epoch 55/100
21/21 [=====] - 8s 406ms/step - loss: 0.159
5 - acc: 0.9429 - val_loss: 0.8862 - val_acc: 0.7600
Epoch 56/100
21/21 [=====] - 8s 403ms/step - loss: 0.128
2 - acc: 0.9417 - val_loss: 1.0566 - val_acc: 0.7800
Epoch 57/100
21/21 [=====] - 8s 398ms/step - loss: 0.165
3 - acc: 0.9466 - val_loss: 0.8299 - val_acc: 0.6600
Epoch 58/100
21/21 [=====] - 8s 397ms/step - loss: 0.186
3 - acc: 0.9320 - val_loss: 1.2668 - val_acc: 0.6000
Epoch 59/100
21/21 [=====] - 8s 400ms/step - loss: 0.162
1 - acc: 0.9417 - val_loss: 0.6520 - val_acc: 0.7000
Epoch 60/100
21/21 [=====] - 8s 405ms/step - loss: 0.109
9 - acc: 0.9515 - val_loss: 0.5125 - val_acc: 0.8200
Epoch 61/100
21/21 [=====] - 8s 399ms/step - loss: 0.152
0 - acc: 0.9466 - val_loss: 0.6334 - val_acc: 0.7200
```

Epoch 62/100
21/21 [=====] - 8s 401ms/step - loss: 0.129
3 - acc: 0.9524 - val_loss: 1.3792 - val_acc: 0.6400
Epoch 63/100
21/21 [=====] - 8s 403ms/step - loss: 0.168
9 - acc: 0.9369 - val_loss: 1.0849 - val_acc: 0.7400
Epoch 64/100
21/21 [=====] - 8s 401ms/step - loss: 0.122
1 - acc: 0.9466 - val_loss: 1.2191 - val_acc: 0.7000
Epoch 65/100
21/21 [=====] - 8s 400ms/step - loss: 0.111
1 - acc: 0.9563 - val_loss: 0.6663 - val_acc: 0.7400
Epoch 66/100
21/21 [=====] - 8s 402ms/step - loss: 0.098
3 - acc: 0.9709 - val_loss: 0.8634 - val_acc: 0.8200
Epoch 67/100
21/21 [=====] - 8s 400ms/step - loss: 0.121
4 - acc: 0.9417 - val_loss: 0.5059 - val_acc: 0.8400
Epoch 68/100
21/21 [=====] - 8s 400ms/step - loss: 0.059
2 - acc: 0.9806 - val_loss: 0.6825 - val_acc: 0.7200
Epoch 69/100
21/21 [=====] - 8s 402ms/step - loss: 0.056
0 - acc: 0.9709 - val_loss: 0.5376 - val_acc: 0.8000
Epoch 70/100
21/21 [=====] - 8s 402ms/step - loss: 0.064
3 - acc: 0.9806 - val_loss: 0.7946 - val_acc: 0.7200
Epoch 71/100
21/21 [=====] - 8s 402ms/step - loss: 0.091
2 - acc: 0.9757 - val_loss: 0.6210 - val_acc: 0.8200
Epoch 72/100
21/21 [=====] - 8s 399ms/step - loss: 0.142
9 - acc: 0.9320 - val_loss: 0.9032 - val_acc: 0.7200
Epoch 73/100
21/21 [=====] - 8s 408ms/step - loss: 0.080
2 - acc: 0.9709 - val_loss: 0.9703 - val_acc: 0.7800
Epoch 74/100
21/21 [=====] - 8s 398ms/step - loss: 0.087
8 - acc: 0.9660 - val_loss: 0.7244 - val_acc: 0.7800
Epoch 75/100
21/21 [=====] - 8s 401ms/step - loss: 0.091
2 - acc: 0.9709 - val_loss: 0.9685 - val_acc: 0.7600
Epoch 76/100
21/21 [=====] - 8s 400ms/step - loss: 0.142
5 - acc: 0.9515 - val_loss: 0.7558 - val_acc: 0.7600
Epoch 77/100
21/21 [=====] - 8s 404ms/step - loss: 0.099
8 - acc: 0.9660 - val_loss: 1.2212 - val_acc: 0.7400
Epoch 78/100
21/21 [=====] - 8s 408ms/step - loss: 0.149
1 - acc: 0.9466 - val_loss: 1.1664 - val_acc: 0.7400
Epoch 79/100
21/21 [=====] - 8s 404ms/step - loss: 0.114
2 - acc: 0.9563 - val_loss: 1.1171 - val_acc: 0.7000
Epoch 80/100
21/21 [=====] - 8s 405ms/step - loss: 0.097
1 - acc: 0.9563 - val_loss: 0.8948 - val_acc: 0.7200
Epoch 81/100
21/21 [=====] - 8s 404ms/step - loss: 0.083
7 - acc: 0.9757 - val_loss: 1.1408 - val_acc: 0.7200
Epoch 82/100

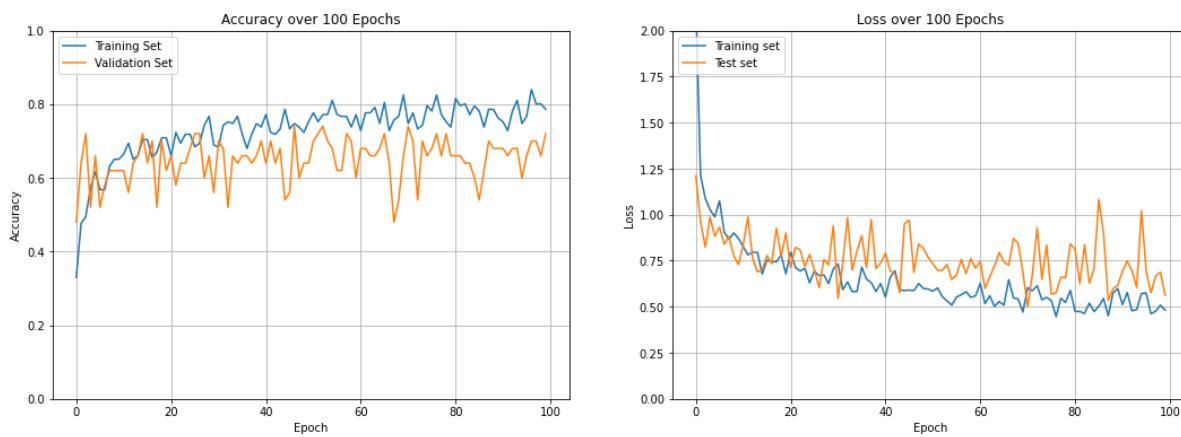
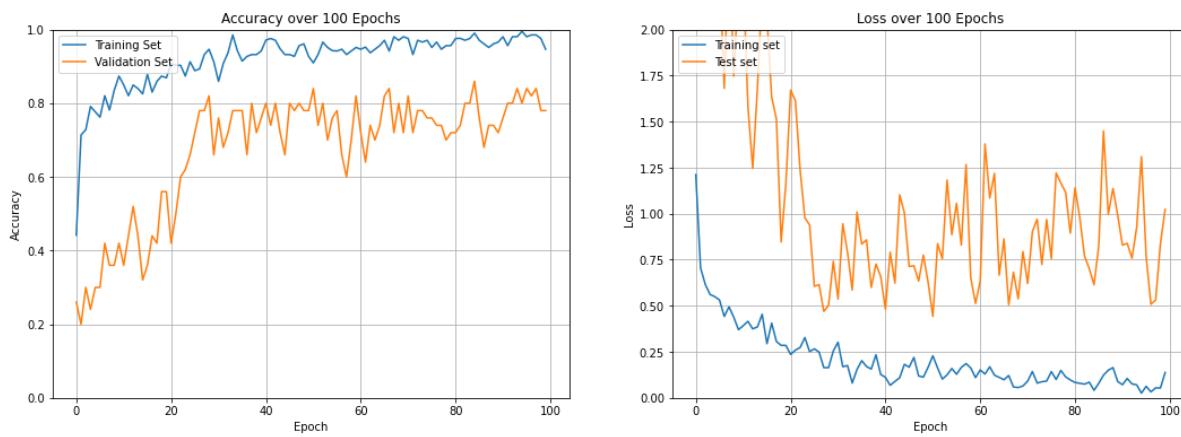
```
21/21 [=====] - 8s 402ms/step - loss: 0.079
1 - acc: 0.9762 - val_loss: 0.9764 - val_acc: 0.7400
Epoch 83/100
21/21 [=====] - 8s 406ms/step - loss: 0.074
6 - acc: 0.9709 - val_loss: 0.7694 - val_acc: 0.8000
Epoch 84/100
21/21 [=====] - 8s 399ms/step - loss: 0.085
2 - acc: 0.9757 - val_loss: 0.7009 - val_acc: 0.8000
Epoch 85/100
21/21 [=====] - 8s 392ms/step - loss: 0.040
8 - acc: 0.9903 - val_loss: 0.6142 - val_acc: 0.8600
Epoch 86/100
21/21 [=====] - 8s 401ms/step - loss: 0.079
4 - acc: 0.9709 - val_loss: 0.8260 - val_acc: 0.7600
Epoch 87/100
21/21 [=====] - 8s 399ms/step - loss: 0.123
5 - acc: 0.9612 - val_loss: 1.4495 - val_acc: 0.6800
Epoch 88/100
21/21 [=====] - 8s 401ms/step - loss: 0.150
0 - acc: 0.9515 - val_loss: 0.9965 - val_acc: 0.7400
Epoch 89/100
21/21 [=====] - 8s 398ms/step - loss: 0.164
4 - acc: 0.9612 - val_loss: 1.1376 - val_acc: 0.7400
Epoch 90/100
21/21 [=====] - 8s 392ms/step - loss: 0.088
5 - acc: 0.9660 - val_loss: 0.9903 - val_acc: 0.7200
Epoch 91/100
21/21 [=====] - 8s 398ms/step - loss: 0.071
4 - acc: 0.9806 - val_loss: 0.8294 - val_acc: 0.7600
Epoch 92/100
21/21 [=====] - 8s 401ms/step - loss: 0.105
1 - acc: 0.9563 - val_loss: 0.8390 - val_acc: 0.8000
Epoch 93/100
21/21 [=====] - 8s 402ms/step - loss: 0.075
9 - acc: 0.9806 - val_loss: 0.7582 - val_acc: 0.8000
Epoch 94/100
21/21 [=====] - 8s 406ms/step - loss: 0.070
6 - acc: 0.9806 - val_loss: 0.9262 - val_acc: 0.8400
Epoch 95/100
21/21 [=====] - 8s 394ms/step - loss: 0.025
6 - acc: 0.9951 - val_loss: 1.3100 - val_acc: 0.8000
Epoch 96/100
21/21 [=====] - 8s 387ms/step - loss: 0.063
0 - acc: 0.9806 - val_loss: 0.7656 - val_acc: 0.8400
Epoch 97/100
21/21 [=====] - 8s 403ms/step - loss: 0.032
5 - acc: 0.9854 - val_loss: 0.5082 - val_acc: 0.8200
Epoch 98/100
21/21 [=====] - 8s 399ms/step - loss: 0.054
4 - acc: 0.9854 - val_loss: 0.5319 - val_acc: 0.8400
Epoch 99/100
21/21 [=====] - 8s 406ms/step - loss: 0.053
4 - acc: 0.9757 - val_loss: 0.8367 - val_acc: 0.7800
Epoch 100/100
21/21 [=====] - 8s 399ms/step - loss: 0.137
7 - acc: 0.9466 - val_loss: 1.0237 - val_acc: 0.7800
```

[5 points] Plot Accuracy and Loss During Training

In [61]:

```
import matplotlib.pyplot as plt
%matplotlib inline

# #Accuracy
# plt.subplot(1,2,1)
def plot_acc_loss(result):
    plt.figure(figsize=(18,6))
    plt.subplot(1,2,1)
    plt.plot(result.history["acc"])
    plt.plot(result.history["val_acc"])
    plt.title("Accuracy over 100 Epochs")
    plt.ylim((0, 1))
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.legend(["Training Set", "Validation Set"], loc="upper left")
    plt.grid(True)
    # Loss
    plt.subplot(1,2,2)
    plt.ylim((0,2))
    plt.plot(result.history["loss"])
    plt.plot(result.history["val_loss"])
    plt.title("Loss over 100 Epochs")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.legend(["Training set", "Test set"], loc="upper left")
    plt.grid(True)
    plt.show()
print("VGG16")
plot_acc_loss(vgg16_model)
print("MobileNet")
plot_acc_loss(mobilenet_model)
```

VGG16**MobileNet****Testing Model**

In [51]:

```
test_datagen = ImageDataGenerator(rescale=1. / 255)

eval_generator = test_datagen.flow_from_directory(TEST_DIR,target_size=IMAGE_SIZE,
                                                 batch_size=1,shuffle=True,seed
                                                 =42,class_mode="categorical")
eval_generator.reset()
print(len(eval_generator))
def Test(model):
    x = model.evaluate_generator(eval_generator,steps = np.ceil(len(eval_generator
))),use_multiprocessing = False,verbose = 1,workers=1)
    return x
```

Found 36 images belonging to 4 classes.

36

In [62]:

```
print('vgg16')
print('Test loss:', Test(vgg16)[0])
print('Test accuracy:',Test(vgg16)[1])
```

```
vgg16
5/36 [==>.....] - ETA: 0s - loss: 1.9903 - acc
c: 0.2000

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:9: User
Warning: `Model.evaluate_generator` is deprecated and will be remove
d in a future version. Please use `Model.evaluate`, which supports g
enerators.
if __name__ == '__main__':
36/36 [=====] - 1s 25ms/step - loss: 0.7133
- acc: 0.7500
Test loss: 0.7132667899131775
36/36 [=====] - 1s 24ms/step - loss: 0.7133
- acc: 0.7500
Test accuracy: 0.75
```

In [28]:

```
print('mobilenet')
print('Test loss:', Test(mobilenet)[0])
print('Test accuracy:', Test(mobilenet)[1])

mobilenet
4/36 [==>.....] - ETA: 0s - loss: 0.7772 - acc: 0.7500

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:9: User
Warning: `Model.evaluate_generator` is deprecated and will be remove
d in a future version. Please use `Model.evaluate`, which supports g
enerators.

if __name__ == '__main__':
    36/36 [=====] - 1s 27ms/step - loss: 1.2211
    - acc: 0.8333
    Test loss: 1.2210861444473267
    36/36 [=====] - 1s 26ms/step - loss: 1.2211
    - acc: 0.8333
    Test accuracy: 0.8333333134651184
```

[10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

In [66]:

```
intermediate_layer_model = models.Model(inputs=vgg16.input,
                                         outputs=vgg16.get_layer('feature_dense')
                                         .output)

tsne_eval_generator = test_datagen.flow_from_directory(DATASET_PATH,target_size=
IMAGE_SIZE,
                                                       batch_size=1,shuffle=False,see
d=42,class_mode="categorical")
tsne_eval_generator.reset()
labels = tsne_eval_generator.classes

X = TSNE().fit_transform(intermediate_layer_model.predict_generator(tsne_eval_ge
nerator, verbose=1))

classes = ["COVID-19", "Normal", "pneumonia_bac", "pneumonia_vir"]
plt.figure(figsize=(18,6))
plt.subplot(1,2,1)
plt.title('VGG16')
for i in range(4):
    cluster = X[np.where(labels == i)]
    plt.scatter(cluster[:, 0], cluster[:, 1], label = classes[i])
plt.legend()

intermediate_layer_model = models.Model(inputs=mobilenet.input,
                                         outputs=mobilenet.get_layer('feature_den
se').output)

tsne_eval_generator = test_datagen.flow_from_directory(DATASET_PATH,target_size=
IMAGE_SIZE,
                                                       batch_size=1,shuffle=False,see
d=42,class_mode="categorical")

labels = tsne_eval_generator.classes

X = TSNE().fit_transform(intermediate_layer_model. predict_generator(tsne_eval_g
enerator, verbose=1))

classes = ["COVID-19", "Normal", "pneumonia_bac", "pneumonia_vir"]
plt.subplot(1,2,2)
plt.title('MobileNet')
for i in range(4):
    cluster = X[np.where(labels == i)]
    plt.scatter(cluster[:, 0], cluster[:, 1], label = classes[i])
plt.legend()
plt.show()
```

Found 270 images belonging to 4 classes.

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:9: User
Warning: `Model.predict_generator` is deprecated and will be removed
in a future version. Please use `Model.predict`, which supports gene
rators.
```

```
if __name__ == '__main__':
```

```
270/270 [=====] - 7s 23ms/step
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:78
3: FutureWarning: The default initialization in TSNE will change fro
m 'random' to 'pca' in 1.2.
```

```
FutureWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:79
3: FutureWarning: The default learning rate in TSNE will change from
200.0 to 'auto' in 1.2.
```

```
FutureWarning,
```

Found 270 images belonging to 4 classes.

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:28: Use
rWarning: `Model.predict_generator` is deprecated and will be remove
d in a future version. Please use `Model.predict`, which supports ge
nerators.
```

```
270/270 [=====] - 7s 24ms/step
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:78
3: FutureWarning: The default initialization in TSNE will change fro
m 'random' to 'pca' in 1.2.
```

```
FutureWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/manifold/_t_sne.py:79
3: FutureWarning: The default learning rate in TSNE will change from
200.0 to 'auto' in 1.2.
```

```
FutureWarning,
```

